

# AI Prompt System Monitor - Technical Specification

## Part 1: Executive Summary, Architecture & Design

**Document Version:** 1.0

**Date:** 2024

**Author:** Yi Zu

**Project:** AI Prompt System Monitor - Production-Ready Full-Stack Application

---

### Table of Contents

1. Executive Summary
  2. Project Overview & Requirements
  3. System Architecture & Design
  4. Component Design
- 

### Executive Summary

#### Business Impact & Value Proposition

The **AI Prompt System Monitor** is a production-ready full-stack application designed to address a critical gap in AI application development: **continuous evaluation and monitoring of prompt systems**. As organizations increasingly rely on LLM-powered applications, the stability and performance of prompt systems become business-critical.

**Target Users:** - **AI/ML Engineers:** Testing prompt variations, monitoring model performance, optimizing costs - **Product Teams:** Ensuring consistent AI behavior across releases, tracking quality metrics - **DevOps/Platform Teams:** Automated monitoring, alerting, and cost management - **Data Scientists:** A/B testing different models, analyzing performance trends

**Key Business Value:** 1. **Cost Control:** Real-time budget tracking prevents unexpected API costs (critical for production deployments) 2. **Quality Assurance:** Automated regression testing ensures prompt changes don't degrade performance 3. **Model Optimization:** Side-by-side comparison enables data-driven model selection 4. **Operational Excellence:** Scheduled monitoring provides early warning of performance degradation 5. **Development Velocity:** Fast feedback loops accelerate prompt engineering iterations

**Market Context:** - LLM applications are moving from prototypes to production systems - Prompt engineering is becoming a core engineering discipline -

Cost overruns from uncontrolled API usage are a common production issue - No existing open-source solution provides this level of integration and automation

### Technical Achievement Summary

This application demonstrates production-grade engineering across the full stack:

- **Backend:** Async-first FastAPI architecture with PostgreSQL, implementing concurrent-safe budget controls, background task processing, and persistent job scheduling
- **Frontend:** Modern React/TypeScript application with real-time progress tracking, performance visualization, and responsive UI
- **AI Integration:** Robust OpenAI API integration with token counting, cost calculation, and error handling
- **Evaluation Engine:** Custom word-overlap similarity algorithm optimized for factual response evaluation
- **Operational Features:** Cron-based scheduling, budget enforcement, cost tracking, and performance alerting

**Completion Status:** - **Level 0:** Core testing infrastructure (prompt management, test execution, evaluation) - **Level 1:** Automated monitoring (scheduling, time-series visualization, alerting) - **Level 2:** Model experimentation (multi-model support, comparison analytics, cost optimization) - **Level 3:** Genetic optimization (not implemented - future enhancement)

---

## Project Overview & Requirements

### Functional Requirements

**Level 0: Core Testing Infrastructure FR-0.1: Prompt System Management - User Story:** As an AI engineer, I need to create and manage prompt templates with variable placeholders so I can test different prompt configurations. - **Acceptance Criteria:** - Create prompt systems with name, description, template, variables, model selection, and model parameters - Support `{variable}` syntax for template interpolation - Store model configuration (temperature, max\_tokens, etc.) as JSONB - CRUD operations via REST API and UI - Validation: template must contain all declared variables - **Implementation:** `backend/app/api/prompt_systems.py`, `backend/app/models/prompt_system.py`

**FR-0.2: Regression Test Management - User Story:** As a QA engineer, I need to define test cases with expected outputs so I can validate prompt system behavior. - **Acceptance Criteria:** - Create regression test sets with input variables and expected output - Associate tests with prompt systems - Support multiple test cases per system - Store input data as

JSONB for flexibility - Delete and update test cases - **Implementation:** backend/app/api/regression\_sets.py

**FR-0.3: Test Execution Engine - User Story:** As a developer, I need to execute regression tests and see results in real-time so I can quickly iterate on prompts. - **Acceptance Criteria:** - Execute all tests for a prompt system - Interpolate template variables with test inputs - Call OpenAI API with configured model and parameters - Display real-time progress (pending → running → completed) - Handle API errors gracefully - Support model override per execution - **Implementation:** backend/app/services/test\_executor.py, backend/app/api/test\_runs.py

**FR-0.4: Evaluation & Scoring - User Story:** As a data scientist, I need similarity scores between actual and expected outputs so I can measure prompt quality. - **Acceptance Criteria:** - Calculate similarity score (0.0-1.0) for each test result - Use word overlap algorithm (Jaccard + recall metrics) - Display scores in results table - Calculate aggregate statistics (average score, success rate) - Store scores in database for historical analysis - **Implementation:** backend/app/services/llm\_service.py::calculate\_similarity()

**FR-0.5: Cost Tracking - User Story:** As a finance manager, I need to see API costs for each test run so I can budget effectively. - **Acceptance Criteria:** - Calculate cost per test based on token usage - Track total cost per test run - Display costs in UI (per test and aggregate) - Store cost data in database - Support cost estimation before execution - **Implementation:** backend/app/services/cost\_calculator.py

**Level 1: Automated Monitoring** **FR-1.1: Scheduled Test Execution - User Story:** As a DevOps engineer, I need to schedule automated test runs so I can monitor prompt systems continuously. - **Acceptance Criteria:** - Create schedules with cron expressions - Support standard cron syntax (minute, hour, day, month, weekday) - Enable/disable schedules without deletion - Execute tests automatically at scheduled times - Track last execution time - Load active schedules on application startup - **Implementation:** backend/app/services/scheduler\_service.py, backend/app/api/schedules.py

**FR-1.2: Performance Trend Visualization - User Story:** As a product manager, I need to see performance trends over time so I can identify degradation early. - **Acceptance Criteria:** - Display time-series chart of average similarity scores - Show test run history with timestamps - Filter by prompt system - Support multiple time ranges (7 days, 30 days, all time) - Interactive chart with tooltips - **Implementation:** frontend/src/components/charts/PerformanceTrendsChart.tsx, frontend/src/hooks/usePerformanceTrends.ts

**FR-1.3: Performance Alerting - User Story:** As an on-call engineer, I need alerts when performance drops so I can investigate issues quickly. - Ac-

**Acceptance Criteria:** - Detect performance degradation (>10% drop in average score) - Display alert banner in UI - Compare current run to historical baseline - Alert persists until next successful run - **Implementation:** frontend/src/components/ui/AlertBanner.tsx

**Level 2: Model Experimentation** **FR-2.1: Multi-Model Support** - **User Story:** As an AI engineer, I need to test the same prompts across different models so I can choose the best model for my use case. - **Acceptance Criteria:** - Support GPT-3.5-turbo, GPT-4o, GPT-4o-mini - Run same test suite on multiple models - Track model used in each test run - Display model name in results - **Implementation:** Model selection in test execution, stored in `test_runs.model_used`

**FR-2.2: Model Comparison Analytics** - **User Story:** As a decision-maker, I need side-by-side model comparisons so I can optimize for accuracy, cost, and speed. - **Acceptance Criteria:** - Compare average similarity scores across models - Compare average response times (latency) - Compare total costs per model - Calculate success rates - Display in tabular format - **Implementation:** Client-side analytics in frontend/src/lib/models/analytics.ts

**FR-2.3: Performance Insights & Recommendations** - **User Story:** As a product owner, I need recommendations on which model to use so I can balance quality and cost. - **Acceptance Criteria:** - Identify “Most Accurate” model (highest average score) - Identify “Fastest” model (lowest average latency) - Identify “Best Value” model (best score-to-cost ratio) - Display recommendations with reasoning - **Implementation:** frontend/src/lib/models/analytics.ts

**FR-2.4: Budget Controls** - **User Story:** As a finance manager, I need budget limits to prevent cost overruns. - **Acceptance Criteria:** - Set daily budget limit per prompt system - Enforce budget before test execution - Auto-pause systems when budget exceeded - Reset daily budget at midnight UTC - Track lifetime and daily costs - Concurrent-safe budget checks (prevent race conditions) - **Implementation:** backend/app/services/budget\_guard.py

## Non-Functional Requirements

**NFR-1: Performance** - Test execution: < 3 seconds per test (GPT-3.5-turbo) - API response time: < 100ms for database queries - UI rendering: < 50ms initial load - Real-time progress updates: 2-second polling interval - Support 100+ concurrent test runs

**NFR-2: Scalability** - Horizontal scaling: Stateless backend services - Database: Connection pooling, indexed queries - Async execution: Non-blocking I/O throughout - Background tasks: Isolated database sessions

**NFR-3: Reliability** - Error handling: Graceful degradation on API failures - Data persistence: All results stored in database - Job recovery: Scheduler

reloads active schedules on restart - Budget safety: Row-level locking prevents race conditions

**NFR-4: Security** - API key management: Environment variables, never logged - Input validation: Pydantic schemas, SQL injection prevention - CORS: Configurable allowed origins - Error messages: No sensitive data in responses

**NFR-5: Usability** - Responsive UI: Works on desktop and tablet - Real-time feedback: Progress bars, status updates - Error messages: Clear, actionable error text - Documentation: Inline help, API docs

### Technical Requirements

**TR-1: Technology Stack** - Backend: Python 3.10+, FastAPI, SQLAlchemy (async), APScheduler - Frontend: React 18, TypeScript, Vite, Tailwind CSS, Recharts - Database: PostgreSQL 12+ (JSONB support required) - AI Provider: OpenAI API (GPT-3.5-turbo, GPT-4o, GPT-4o-mini)

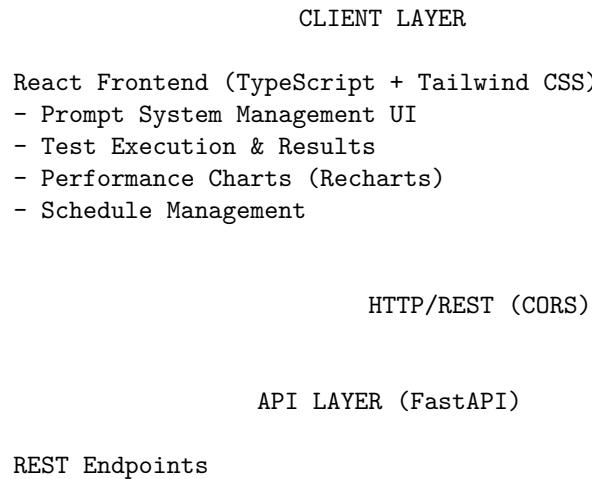
**TR-2: Integration Requirements** - OpenAI API: Async HTTP client, token counting, error handling - Database: Async SQLAlchemy with connection pooling - Scheduler: APScheduler with UTC timezone, persistent jobs

**TR-3: Deployment Requirements** - Development: Local PostgreSQL (Supabase recommended) - Production: Cloud-ready (Docker, Kubernetes compatible) - Environment: Configurable via environment variables - Monitoring: Health check endpoints

---

## System Architecture & Design

### High-Level Architecture



- /api/prompt-systems (CRUD)
- /api/regression-sets (CRUD)
- /api/test-runs (Create, List, Get)
- /api/schedules (CRUD, Trigger)

#### Background Tasks (FastAPI BackgroundTasks)

- Async test execution
- Isolated database sessions

### SERVICE LAYER

LLMService	TestExecutor	Scheduler
- Template interp.	- Orchestrate execution	- APScheduler
- OpenAI API	- Progress tracking	- Cron jobs
- Evaluation		- Job mgmt

CostCalculator	BudgetGuard
- Token calc	- Budget
- Pricing	enforcement

### DATA LAYER (PostgreSQL)

prompt_systems	regression_sets	test_runs
----------------	-----------------	-----------

test_results	schedules
--------------	-----------

### EXTERNAL SERVICES

OpenAI API
- Chat Completions
- Token Usage Reporting

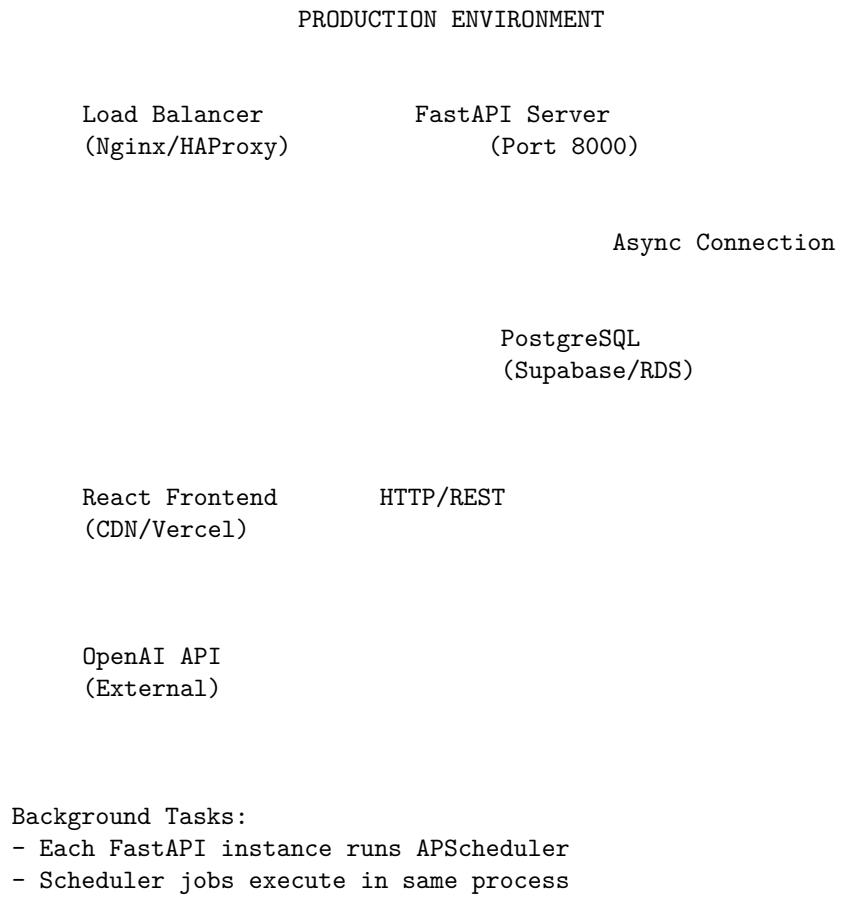
## Data Flow: Test Execution

1. User clicks "Run Tests" in UI
2. Frontend: POST /api/test-runs
  - {
    - "system\_id": "uuid",
    - "model\_override": "gpt-4o"
  - }
3. API: TestExecutor.prepare\_test\_run()
  - Load prompt system from DB
  - Load regression tests
  - Estimate cost
  - BudgetGuard.check\_and\_reserve\_budget() [WITH FOR UPDATE lock]
  - Create TestRun record (status: "pending")
  - Return test\_run\_id immediately
4. API: Background task starts
  - TestExecutor.execute\_test\_run\_by\_id()
  - Update status: "running"
5. For each RegressionSet:
  - LLMService.interpolate\_template()
    - Replace {variables} with test inputs
  - LLMService.execute\_prompt()
    - Call OpenAI API (async)
    - Measure execution time
    - Extract token usage
  - CostCalculator.calculate\_cost()
    - Compute USD cost from tokens
  - LLMService.calculate\_similarity()
    - Word overlap algorithm
    - Jaccard + recall metrics
  - Create TestResult record

- Store actual\_output, score, cost, tokens
- Update TestRun.completed\_tests++
  

  6. TestExecutor.\_finalize\_test\_run()
    - Calculate avg\_score
    - Sum total\_cost, total\_tokens
    - BudgetGuard.record\_actual\_cost() [adjust for estimation]
    - Update status: "completed"
  
  7. Frontend: Poll GET /api/test-runs/{id} every 2s
    - Display progress: completed\_tests / total\_tests
    - Show results table when status = "completed"

### Deployment Topology



- For high availability: Use external job queue (Redis + Celery)

### Service Boundaries & Responsibilities

**API Layer (backend/app/api/)** - **Responsibility:** HTTP request/response handling, validation, routing - **Pattern:** RESTful endpoints, dependency injection for DB sessions - **Error Handling:** HTTPException with appropriate status codes - **Security:** CORS middleware, input validation via Pydantic

**Service Layer (backend/app/services/)** - **LLMService:** OpenAI integration, template interpolation, evaluation - **TestExecutor:** Test orchestration, progress tracking, result aggregation - **SchedulerService:** APScheduler management, cron job lifecycle - **CostCalculator:** Token-to-cost conversion, pricing management - **BudgetGuard:** Concurrent-safe budget enforcement

**Data Layer (backend/app/models/, backend/app/schemas/)** - **Models:** SQLAlchemy ORM definitions, database relationships - **Schemas:** Pydantic models for request/response validation - **Database:** Async SQLAlchemy with connection pooling

**Frontend (frontend/src/)** - **Components:** React components (atomic design pattern) - **Services:** API client with Axios, error handling - **State:** React Query for server state, local state for UI - **Visualization:** Recharts for performance trends

---

## Component Design

### Backend Architecture

#### API Layer (FastAPI) Route Design Patterns:

```
# Standard REST pattern with async/await
@router.post("/prompt-systems", response_model=PromptSystemResponse, status_code=201)
async def create_prompt_system(
    system: PromptSystemCreate,
    db: AsyncSession = Depends(get_db)
):
    # Validation handled by Pydantic
    db_system = PromptSystem(**system.model_dump())
    db.add(db_system)
    await db.commit()
    await db.refresh(db_system)
    return db_system
```

#### Request/Response Schemas (Pydantic):

```
# Request Schema
class PromptSystemCreate(BaseModel):
```

```

name: str = Field(..., min_length=1, max_length=255)
description: Optional[str] = None
template: str = Field(..., min_length=1)
variables: List[str] = Field(default_factory=list)
model: str = Field(default="gpt-3.5-turbo")
model_params: Dict[str, Any] = Field(default_factory=dict)

# Response Schema
class PromptSystemResponse(BaseModel):
    id: UUID
    name: str
    template: str
    variables: List[str]
    model: str
    model_params: Dict[str, Any]
    created_at: datetime

    class Config:
        from_attributes = True # For SQLAlchemy models

```

#### Error Handling Strategy:

```

# Centralized error handling
try:
    executor = TestExecutor(db)
    test_run_id = await executor.prepare_test_run(...)
except ValueError as e:
    raise HTTPException(status_code=400, detail=str(e))
except BudgetExceededException as e:
    raise HTTPException(status_code=402, detail=str(e)) # Payment Required
except Exception as e:
    logger.error(f"Unexpected error: {e}", exc_info=True)
    raise HTTPException(status_code=500, detail="Internal server error")

```

#### CORS Configuration:

```

# backend/app/main.py
app.add_middleware(
    CORSMiddleware,
    allow_origins=settings.cors_origins, # From environment
    allow_credentials=True,
    allow_methods=["*"],
    allow_headers=["*"],
)

```

#### Service Layer LLMService: OpenAI Integration

```

class LLMService:
    def __init__(self):
        self.client = AsyncOpenAI(api_key=settings.OPENAI_API_KEY)
        self.cost_calculator = CostCalculator()

    @async def execute_prompt(
        self,
        template: str,
        variables: Dict[str, Any],
        model: str = "gpt-4o",
        model_params: Dict[str, Any] = None
    ) -> ExecutionResult:
        # 1. Interpolate template
        prompt = self.interpolate_template(template, variables)

        # 2. Prepare parameters with defaults
        params = model_params or {}
        params.setdefault('temperature', 0.7)
        params.setdefault('max_tokens', 1000)

        # 3. Execute with timing
        start_time = time.time()
        response = await self.client.chat.completions.create(
            model=model,
            messages=[{"role": "user", "content": prompt}],
            **params
        )
        execution_time = int((time.time() - start_time) * 1000)

        # 4. Extract metrics
        usage = response.usage
        cost = self.cost_calculator.calculate_cost(
            model=model,
            input_tokens=usage.prompt_tokens,
            output_tokens=usage.completion_tokens
        )

    return ExecutionResult(
        output=response.choices[0].message.content,
        execution_time_ms=execution_time,
        cost_usd=cost,
        total_tokens=usage.total_tokens
    )

```

**Token Counting & Rate Limiting:** - OpenAI API automatically returns token usage in response - No client-side token counting needed (more accurate)

- Rate limiting handled by OpenAI (429 responses caught and logged) - Future: Implement exponential backoff retry logic

### TestExecutor: Async Execution Patterns

```
class TestExecutor:
    async def execute_test_run_by_id(self, test_run_id: UUID, ...):
        # Load test run
        test_run = await self._load_test_run(test_run_id)
        test_run.status = "running"
        await self.db.commit()

        # Execute tests sequentially (can be parallelized)
        for regression_set in regression_sets:
            test_result = await self._execute_single_test(...)
            test_run.completed_tests += 1
            await self.db.commit() # Progress visible immediately

        # Finalize
        await self._finalize_test_run(test_run, ...)
```

### Background Task Management:

```
# API endpoint returns immediately
@router.post("/test-runs", status_code=202)
async def create_test_run(..., background_tasks: BackgroundTasks):
    test_run_id = await executor.prepare_test_run(...)

    # Execute in background with isolated session
    background_tasks.add_task(
        _execute_test_run_background,
        test_run_id,
        model_override
    )

    return {"test_run_id": str(test_run_id), "status": "pending"}

# Background task creates its own DB session
async def _execute_test_run_background(test_run_id: UUID, ...):
    async with AsyncSessionLocal() as db:
        executor = TestExecutor(db)
        await executor.execute_test_run_by_id(test_run_id, ...)
```

### SchedulerService: APScheduler Configuration

```
class SchedulerService:
    def __init__(self):
        # Singleton pattern
        self.scheduler = AsyncIOScheduler(timezone="UTC")
```

```

    async def start(self):
        self.scheduler.start()
        await self._load_schedules_from_db() # Restore on startup

    async def _add_job(self, schedule_id: UUID, system_id: UUID, cron: str):
        trigger = CronTrigger.from_crontab(cron, timezone="UTC")
        self.scheduler.add_job(
            func=self._execute_scheduled_test,
            trigger=trigger,
            id=str(schedule_id),
            args=[system_id, schedule_id],
            replace_existing=True,
            misfire_grace_time=300 # 5 min grace period
        )

```

**Job Persistence:** - Current: In-memory (lost on restart, reloaded from DB) -  
 Future: Use APScheduler's SQLAlchemy job store for true persistence

#### Cost Tracking Implementation:

```

class CostCalculator:
    MODEL_PRICING = {
        "gpt-4o": {"input": 2.50, "output": 10.00},
        "gpt-4o-mini": {"input": 0.150, "output": 0.600},
        "gpt-3.5-turbo": {"input": 0.50, "output": 1.50}
    }

    @classmethod
    def calculate_cost(cls, model: str, input_tokens: int, output_tokens: int):
        pricing = cls.MODEL_PRICING.get(model, cls.MODEL_PRICING["gpt-4o"])
        input_cost = (input_tokens / 1_000_000) * pricing["input"]
        output_cost = (output_tokens / 1_000_000) * pricing["output"]
        return input_cost + output_cost

    @classmethod
    def estimate_cost(cls, model: str, num_tests: int):
        # Conservative estimate with 20% buffer
        avg_input = 500
        avg_output = 200
        cost_per_test = cls.calculate_cost(model, avg_input, avg_output)
        return cost_per_test * num_tests * 1.2 # 20% safety margin

```

#### Data Layer Database Schema with Relationships:

```

-- Core entities
CREATE TABLE prompt_systems (

```

```

    id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
    name VARCHAR(255) NOT NULL,
    template TEXT NOT NULL,
    variables JSONB NOT NULL DEFAULT '[]',
    model VARCHAR(100) NOT NULL DEFAULT 'gpt-3.5-turbo',
    model_params JSONB NOT NULL DEFAULT '{}',
    -- Budget controls
    max_daily_cost_usd FLOAT NOT NULL DEFAULT 1.0,
    daily_cost_usd FLOAT NOT NULL DEFAULT 0.0,
    total_cost_usd FLOAT NOT NULL DEFAULT 0.0,
    last_reset_date DATE,
    is_budget_paused BOOLEAN NOT NULL DEFAULT FALSE,
    created_at TIMESTAMP WITH TIME ZONE DEFAULT NOW()
);

CREATE TABLE regression_sets (
    id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
    system_id UUID NOT NULL REFERENCES prompt_systems(id) ON DELETE CASCADE,
    inputs JSONB NOT NULL,
    expected_output TEXT NOT NULL,
    created_at TIMESTAMP WITH TIME ZONE DEFAULT NOW()
);

CREATE TABLE test_runs (
    id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
    system_id UUID NOT NULL REFERENCES prompt_systems(id),
    status VARCHAR(50) NOT NULL DEFAULT 'pending',
    model_used VARCHAR(100) NOT NULL,
    triggered_by VARCHAR(50) DEFAULT 'manual',
    total_tests INTEGER DEFAULT 0,
    completed_tests INTEGER DEFAULT 0,
    avg_score DECIMAL(4,3),
    started_at TIMESTAMP WITH TIME ZONE DEFAULT NOW(),
    completed_at TIMESTAMP WITH TIME ZONE,
    estimated_cost_usd FLOAT,
    actual_cost_usd FLOAT,
    total_tokens_used INTEGER
);

CREATE TABLE test_results (
    id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
    test_run_id UUID NOT NULL REFERENCES test_runs(id) ON DELETE CASCADE,
    regression_id UUID NOT NULL REFERENCES regression_sets(id),
    actual_output TEXT NOT NULL,
    expected_output TEXT NOT NULL,
    score DECIMAL(4,3) NOT NULL,

```

```

        execution_time_ms INTEGER NOT NULL,
        cost_usd FLOAT,
        tokens_used INTEGER,
        created_at TIMESTAMP WITH TIME ZONE DEFAULT NOW()
    );

CREATE TABLE schedules (
    id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
    system_id UUID NOT NULL REFERENCES prompt_systems(id) ON DELETE CASCADE,
    name VARCHAR(255) NOT NULL,
    cron_expression VARCHAR(50) NOT NULL,
    active BOOLEAN DEFAULT true,
    last_run_at TIMESTAMP WITH TIME ZONE,
    created_at TIMESTAMP WITH TIME ZONE DEFAULT NOW()
);

-- Performance indexes
CREATE INDEX idx_test_runs_system_id ON test_runs(system_id);
CREATE INDEX idx_test_results_test_run_id ON test_results(test_run_id);
CREATE INDEX idx_prompt_systems_last_reset_date ON prompt_systems(last_reset_date);

```

#### Connection Pooling Configuration:

```

# backend/app/core/database.py
engine = create_async_engine(
    settings.DATABASE_URL,
    echo=True, # Log SQL (disable in production)
    pool_pre_ping=True, # Verify connections
    pool_size=10, # Default connection pool
    max_overflow=20 # Additional connections under load
)

```

#### Query Optimization Patterns:

```

# Eager loading to avoid N+1 queries
result = await db.execute(
    select(TestRun)
    .where(TestRun.system_id == system_id)
    .order_by(TestRun.started_at.desc())
    .limit(limit)
    .options(selectinload(TestRun.test_results)) # Load related results
)

```

#### Frontend Architecture

##### Component Hierarchy

App.tsx

```

PromptSystemList (Page)
  PromptSystemCard (Feature)
    Badge (UI)
    Button (UI)
  PromptSystemForm (Feature - Modal)
    Input (UI)
    Button (UI)
PromptSystemDetail (Page)
  RegressionSetList (Feature)
    TestResultCard (Feature)
  TestRunsView (Feature)
    PerformanceTrendsChart (Feature)
    Recharts components
  BudgetControl (UI)
ScheduleList (Page)
  ScheduleCard (Feature)
  ScheduleForm (Feature - Modal)

```

**Atomic Design Implementation:** - **UI Components:** Reusable primitives (Button, Input, Card, Badge, Modal) - **Feature Components:** Domain-specific (PromptSystemForm, TestRunsView) - **Page Components:** Top-level views (PromptSystemList, PromptSystemDetail)

### State Management Patterns

```

// Server state: React Query
const { data: systems, isLoading } = useQuery({
  queryKey: ['prompt-systems'],
  queryFn: () => promptSystemsApi.getAll()
});

// Local state: useState for UI
const [isModalOpen, setIsModalOpen] = useState(false);

// Optimistic updates
const mutation = useMutation({
  mutationFn: promptSystemsApi.create,
  onMutate: async (newSystem) => {
    // Cancel outgoing refetches
    await queryClient.cancelQueries(['prompt-systems']);
    // Snapshot previous value
    const previous = queryClient.getQueryData(['prompt-systems']);
    // Optimistically update
    queryClient.setQueryData(['prompt-systems'], (old) => [...old, newSystem]);
    return { previous };
},

```

```

        onError: (err, newSystem, context) => {
          // Rollback on error
          queryClient.setQueryData(['prompt-systems'], context.previous);
        }
      );
    );
  
```

## API Integration Layer

```

// frontend/src/services/api.ts
export const apiClient = axios.create({
  baseURL: import.meta.env.VITE_API_URL || 'http://localhost:8000',
  headers: { 'Content-Type': 'application/json' }
});

// Request interceptor (for auth tokens)
apiClient.interceptors.request.use((config) => {
  // Add auth token if available
  return config;
});

// Response interceptor (error handling)
apiClient.interceptors.response.use(
  (response) => response,
  (error) => {
    console.error('API Error:', error.response?.data || error.message);
    return Promise.reject(error);
  }
);

```

**UI/UX Design** - **Design System:** - **Colors:** Tailwind CSS default palette, custom accent colors - **Typography:** System fonts (San Francisco on macOS, Segoe UI on Windows) - **Spacing:** Consistent 4px grid (Tailwind spacing scale) - **Components:** Consistent border radius, shadows, hover states

**Responsive Design:** - Mobile-first approach - Breakpoints: sm (640px), md (768px), lg (1024px), xl (1280px) - Grid layout adapts to screen size - Tables scroll horizontally on mobile

**Accessibility:** - Semantic HTML (buttons, inputs, labels) - ARIA labels for icons - Keyboard navigation support - Focus indicators visible

**Performance Optimization:** - Code splitting: React.lazy() for route-based splitting - Image optimization: Vite handles asset optimization - Bundle size: Tree-shaking removes unused code - Memoization: React.memo() for expensive components

---

**End of Part 1**

*Continue to Part 2 for: Detailed Implementation, Technical Deep Dives, Performance & Scalability*

# AI Prompt System Monitor - Technical Specification

## Part 2: Detailed Implementation & Technical Deep Dives

**Document Version:** 1.0

**Date:** 2024

**Author:** Yi Zu

**Project:** AI Prompt System Monitor - Production-Ready Full-Stack Application

---

### Table of Contents

1. Detailed Implementation
  2. Technical Deep Dives
  3. Performance & Scalability
  4. Advanced Features
- 

## Detailed Implementation

### Core Business Logic

#### Prompt System Management Template Interpolation Implementation:

```
# backend/app/services/llm_service.py
@staticmethod
def interpolate_template(template: str, variables: Dict[str, Any]) -> str:
    """
    Replace {variable} placeholders with actual values.
    """
```

*Simple string replacement - no advanced templating engine needed.  
Handles edge cases: missing variables, empty values, type conversion.*

*Args:*

*template: String with {variable} placeholders  
variables: Dictionary of variable names to values*

*Returns:*

*Interpolated string with variables replaced*

*Example:*

```
>>> template = "Hello {name}, you have {count} messages."
>>> variables = {"name": "Alice", "count": 5}
```

```

>>> interpolate_template(template, variables)
"Hello Alice, you have 5 messages."
"""

result = template
for key, value in variables.items():
    placeholder = f"{{{{key}}}}" # {{variable_name}}
    result = result.replace(placeholder, str(value))
return result

```

#### Validation Strategy:

```

# backend/app/schemas/prompt_system.py
class PromptSystemCreate(BaseModel):
    name: str = Field(..., min_length=1, max_length=255)
    template: str = Field(..., min_length=1)
    variables: List[str] = Field(default_factory=list)

    @validator('template')
    def validate_template_has_variables(cls, v, values):
        """Ensure all declared variables exist in template."""
        variables = values.get('variables', [])
        for var in variables:
            if f"{{{{var}}}}" not in v:
                raise ValueError(f"Variable '{var}' not found in template")
        return v

```

#### Variable Extraction (Future Enhancement):

```

# Could auto-detect variables from template
import re

def extract_variables(template: str) -> List[str]:
    """Extract {{variable}} placeholders from template."""
    pattern = r'\{{(\w+)\}\}'
    return list(set(re.findall(pattern, template)))

```

#### Test Execution Engine Async Execution Patterns:

```

# backend/app/services/test_executor.py
async def execute_test_run_by_id(self, test_run_id: UUID, ...):
    """
    Execute tests for an existing test run.

```

*This method is designed for background execution:*

- Updates status incrementally (visible to polling clients)
- Handles errors gracefully (creates failed test results)
- Tracks progress (completed\_tests / total\_tests)
- Finalizes with aggregate statistics

```

"""
test_run = await self._load_test_run(test_run_id)
test_run.status = "running"
await self.db.commit()  # Make status visible immediately

# Execute tests sequentially
# Note: Could parallelize with asyncio.gather() for speed
total_score = 0.0
total_cost = 0.0
successful_tests = 0

for regression_set in regression_sets:
    try:
        test_result = await self._execute_single_test(
            test_run=test_run,
            system=system,
            regression_set=regression_set,
            model_override=model_override
        )

        # Aggregate metrics
        total_score += float(test_result.score)
        total_cost += test_result.cost_usd
        successful_tests += 1
        test_run.completed_tests += 1

        # Commit after each test (progress visible)
        await self.db.commit()

    except Exception as e:
        logger.error(f"Test failed: {e}")
        await self._create_failed_test_result(...)
        test_run.completed_tests += 1
        await self.db.commit()

    # Finalize with aggregate statistics
    await self._finalize_test_run(test_run, total_score, successful_tests, ...)

Parallel Execution (Future Enhancement):
```

```

# Could execute tests in parallel for speed
async def execute_test_run_parallel(self, ...):
    tasks = [
        self._execute_single_test(...)
        for regression_set in regression_sets
    ]
    results = await asyncio.gather(*tasks, return_exceptions=True)
```

```

# Handle results and exceptions
for i, result in enumerate(results):
    if isinstance(result, Exception):
        await self._create_failed_test_result(...)
    else:
        # Process successful result
    ...

```

#### Background Task Management:

```

# backend/app/api/test_runs.py
async def _execute_test_run_background(test_run_id: UUID, model_override: str = None):
    """
    Background task with isolated database session.

    Critical: Creates its own session to avoid:
    - Session lifecycle issues (FastAPI dependency sessions close after request)
    - Transaction conflicts
    - Connection pool exhaustion
    """

    async with AsyncSessionLocal() as db:
        try:
            executor = TestExecutor(db)
            await executor.execute_test_run_by_id(test_run_id, model_override)
        except Exception as e:
            logger.error(f"Background execution failed: {e}")
            # Update test run status to failed
            result = await db.execute(
                select(TestRun).where(TestRun.id == test_run_id)
            )
            test_run = result.scalar_one_or_none()
            if test_run:
                test_run.status = "failed"
                await db.commit()

```

#### Progress Tracking Implementation:

```

# Frontend polling pattern
const pollTestRun = async (testRunId: string) => {
    const pollInterval = 2000; // 2 seconds
    const maxAttempts = 300; // 10 minutes max

    for (let i = 0; i < maxAttempts; i++) {
        const testRun = await testRunsApi.getById(testRunId);

        if (testRun.status === 'completed' || testRun.status === 'failed') {
            return testRun;
        }
    }
}

```

```

    }

    // Update progress bar
    setProgress(testRun.completed_tests / testRun.total_tests);

    await new Promise(resolve => setTimeout(resolve, pollInterval));
}

throw new Error('Test run timeout');
};

```

#### Error Recovery Mechanisms:

```

# Retry logic for transient failures
async def execute_prompt_with_retry(
    self,
    template: str,
    variables: Dict[str, Any],
    model: str,
    max_retries: int = 3
) -> ExecutionResult:
    """Execute with exponential backoff retry."""
    for attempt in range(max_retries):
        try:
            return await self.execute_prompt(template, variables, model)
        except Exception as e:
            if attempt == max_retries - 1:
                raise # Final attempt failed

    # Exponential backoff: 1s, 2s, 4s
    wait_time = 2 ** attempt
    logger.warning(f'Retry {attempt + 1}/{max_retries} after {wait_time}s')
    await asyncio.sleep(wait_time)

```

#### Evaluation Metrics Word Overlap Algorithm - Detailed Implementation:

```

# backend/app/services/llm_service.py
@staticmethod
def calculate_similarity(actual: str, expected: str) -> float:
    """
    Calculate similarity score using word overlap approach.

    Algorithm:
    1. Normalize strings (lowercase, strip whitespace)
    2. Check exact match (score = 1.0)
    3. Check substring match (score = 0.95)

```

4. Calculate Jaccard similarity (*intersection / union*)
5. Calculate recall (*intersection / expected\_words*)
6. Weighted average: Jaccard (30%) + Recall (70%)

*Why this approach:*

- *Fast:  $O(n)$  where  $n$  is word count*
- *Cost-free: No API calls*
- *Effective: Good for factual responses with key terms*
- *Interpretable: Scores correlate with human judgment*

*Trade-offs:*

- *Less sophisticated than embeddings (but no cost/latency)*
- *May miss semantic similarity (e.g., "happy" vs "joyful")*
- *Works best when expected output contains key terms*

*Returns:*

```

        float: Similarity score 0.0-1.0
"""

# Step 1: Normalize
if not expected or not actual:
    return 0.0

actual_lower = actual.lower().strip()
expected_lower = expected.lower().strip()

# Step 2: Exact match
if actual_lower == expected_lower:
    return 1.0

# Step 3: Substring match (expected contained in actual)
if expected_lower in actual_lower:
    return 0.95 # High score but not perfect (distinguish from exact)

# Step 4: Word-based comparison
actual_words = set(actual_lower.split())
expected_words = set(expected_lower.split())

if not expected_words:
    return 0.0

# Step 5: Calculate Jaccard similarity
intersection = len(actual_words & expected_words)
union = len(actual_words | expected_words)

if union == 0:
    return 0.0

```

```

jaccard_score = intersection / union

# Step 6: Calculate recall (coverage of expected words)
recall = intersection / len(expected_words)

# Step 7: Weighted average
# Recall weighted higher (70%) because we care more about
# covering expected content than avoiding extra words
score = (jaccard_score * 0.3) + (recall * 0.7)

return min(score, 1.0) # Cap at 1.0

```

#### Example Calculations:

##### Example 1: Exact Match

actual: "We are open Monday-Friday 9am-5pm PT"  
expected: "We are open Monday-Friday 9am-5pm PT"  
→ Score: 1.0 (exact match)

##### Example 2: Substring Match

actual: "Our business hours are: We are open Monday-Friday 9am-5pm PT. Thank you!"  
expected: "We are open Monday-Friday 9am-5pm PT"  
→ Score: 0.95 (substring match)

##### Example 3: Word Overlap

actual: "We operate Monday through Friday, 9am to 5pm Pacific Time"  
expected: "We are open Monday-Friday 9am-5pm PT"

actual\_words: {we, operate, monday, through, friday, 9am, to, 5pm, pacific, time}  
expected\_words: {we, are, open, monday-friday, 9am-5pm, pt}  
intersection: {we} (only "we" matches exactly)  
union: {we, operate, monday, through, friday, 9am, to, 5pm, pacific, time, are, open, monda...}  
  
jaccard = 1/14 = 0.071  
recall = 1/6 = 0.167  
score = (0.071 \* 0.3) + (0.167 \* 0.7) = 0.135  
  
→ Score: 0.135 (low due to different phrasing)

##### Example 4: Good Overlap

actual: "We are open Monday-Friday from 9am to 5pm Pacific Time"  
expected: "We are open Monday-Friday 9am-5pm PT"

actual\_words: {we, are, open, monday-friday, from, 9am, to, 5pm, pacific, time}  
expected\_words: {we, are, open, monday-friday, 9am-5pm, pt}  
intersection: {we, are, open, monday-friday}

```

union: {we, are, open, monday-friday, from, 9am, to, 5pm, pacific, time, 9am-5pm, pt}

jaccard = 4/12 = 0.333
recall = 4/6 = 0.667
score = (0.333 * 0.3) + (0.667 * 0.7) = 0.567

→ Score: 0.567 (good overlap)

```

#### Alternative Evaluation Approaches:

##### 1. BLEU Score (Machine Translation Metric):

```

from nltk.translate.bleu_score import sentence_bleu

def calculate_bleu(actual: str, expected: str) -> float:
    """BLEU score for n-gram overlap."""
    actual_tokens = actual.lower().split()
    expected_tokens = expected.lower().split()
    reference = [expected_tokens] # BLEU expects list of references
    return sentence_bleu(reference, actual_tokens)

```

**Pros:** Standard metric, handles n-grams

**Cons:** Designed for translation, less intuitive scores

##### 2. ROUGE Score (Summarization Metric):

```

from rouge_score import rouge_scorer

def calculate_rouge(actual: str, expected: str) -> float:
    """ROUGE-L (longest common subsequence)."""
    scorer = rouge_scorer.RougeScorer(['rougeL'], use_stemmer=True)
    scores = scorer.score(expected, actual)
    return scores['rougeL'].fmeasure

```

**Pros:** Good for summarization tasks

**Cons:** Requires additional dependency, slower

##### 3. Embedding-Based Similarity:

```

from sentence_transformers import SentenceTransformer

model = SentenceTransformer('all-MiniLM-L6-v2')

def calculate_embedding_similarity(actual: str, expected: str) -> float:
    """Cosine similarity of sentence embeddings."""
    embeddings = model.encode([actual, expected])
    similarity = cosine_similarity([embeddings[0]], [embeddings[1]])[0][0]
    return float(similarity)

```

**Pros:** Captures semantic similarity

**Cons:** Requires model download, slower, memory usage

#### 4. LLM-as-Judge:

```
async def llm_judge_similarity(actual: str, expected: str) -> float:
    """Use LLM to evaluate similarity."""
    prompt = f"""
        Rate the similarity between these two responses on a scale of 0.0 to 1.0.

        Expected: {expected}
        Actual: {actual}

        Respond with only a number between 0.0 and 1.0.
    """
    response = await llm_service.execute_prompt(prompt, {})
    score = float(response.output.strip())
    return score
```

**Pros:** Most sophisticated, understands context

**Cons:** Expensive (extra API call per test), slower, variable quality

#### 5. Custom Evaluation Functions (Future):

```
# Allow users to upload Python evaluation functions
def custom_evaluator(actual: str, expected: str, metadata: dict) -> float:
    """User-defined evaluation logic."""
    # Could check for specific keywords, regex patterns, etc.
    if "error" in actual.lower():
        return 0.0
    # ... custom logic
    return calculate_similarity(actual, expected)
```

**Current Choice Rationale:** - Word overlap chosen for MVP because: - Zero cost (no API calls) - Fast execution (< 1ms) - Good enough for factual responses - Easy to understand and debug - Can be enhanced later with embeddings or LLM-as-judge

### Advanced Features

#### Multi-Model Support Model Abstraction Layer:

```
# Current: Direct OpenAI integration
# Future: Provider abstraction

class LLMPromoter(ABC):
    @abstractmethod
    async def complete(self, prompt: str, model: str, **params) -> CompletionResult:
        pass

class OpenAIProvider(LLMPromoter):
    async def complete(self, prompt: str, model: str, **params):
```

```

# Current implementation
...
# Future: Anthropic Claude support
...

class AnthropicProvider(LLMProvider):
    async def complete(self, prompt: str, model: str, **params):
        ...
        ...

class LLMService:
    def __init__(self, provider: LLMProvider):
        self.provider = provider

    async def execute_prompt(self, ...):
        return await self.provider.complete(...)
```

#### Provider-Specific Configurations:

```

# Model capabilities and limits
MODEL_CONFIGS = {
    "gpt-4o": {
        "max_tokens": 16384,
        "supports_streaming": True,
        "supports_function_calling": True,
        "cost_per_1k_input": 0.0025,
        "cost_per_1k_output": 0.01
    },
    "gpt-4o-mini": {
        "max_tokens": 16384,
        "supports_streaming": True,
        "supports_function_calling": True,
        "cost_per_1k_input": 0.00015,
        "cost_per_1k_output": 0.0006
    },
    "gpt-3.5-turbo": {
        "max_tokens": 16384,
        "supports_streaming": True,
        "supports_function_calling": True,
        "cost_per_1k_input": 0.0005,
        "cost_per_1k_output": 0.0015
    }
}
```

#### Fallback Strategies:

```

async def execute_prompt_with_fallback(
    self,
    template: str,
    variables: Dict[str, Any],
```

```

    primary_model: str,
    fallback_model: str = "gpt-3.5-turbo"
) -> ExecutionResult:
    """Execute with automatic fallback on failure."""
    try:
        return await self.execute_prompt(template, variables, primary_model)
    except Exception as e:
        logger.warning(f"Primary model {primary_model} failed: {e}, trying fallback")
        return await self.execute_prompt(template, variables, fallback_model)

```

### Cost Optimization Across Models:

```

def recommend_model(
    self,
    required_quality: float = 0.8,
    max_cost_per_test: float = 0.01
) -> str:
    """Recommend model based on quality and cost constraints."""
    # Analyze historical data
    model_performance = self.get_model_performance_stats()

    # Filter by quality requirement
    candidates = [
        m for m in model_performance
        if m['avg_score'] >= required_quality
    ]

    # Filter by cost constraint
    candidates = [
        m for m in candidates
        if m['avg_cost'] <= max_cost_per_test
    ]

    # Return best value (score/cost ratio)
    if candidates:
        return max(candidates, key=lambda m: m['avg_score'] / m['avg_cost'])['model']
    else:
        return "gpt-3.5-turbo" # Default fallback

```

### Scheduling & Automation Cron Expression Handling:

```

# backend/app/services/scheduler_service.py
from apscheduler.triggers.cron import CronTrigger

def parse_cron_expression(cron: str) -> CronTrigger:
    """
    Parse cron expression with validation.

```

```

Format: "minute hour day month weekday"
Examples:
    "0 * * * *"      - Every hour at minute 0
    "0 9 * * 1-5"    - 9 AM on weekdays
    "*/15 * * * *"   - Every 15 minutes
    "0 0 1 * *"     - First day of month at midnight
    """
try:
    trigger = CronTrigger.from_crontab(cron, timezone="UTC")
    return trigger
except ValueError as e:
    raise ValueError(f"Invalid cron expression: {cron}. Error: {e}")

```

#### Cron Expression Helper (Frontend):

```

// frontend/src/components/schedules/CronExpressionHelper.tsx
export const CRON_PRESETS = {
    'Every Minute': '*/1 * * * *',
    'Every 5 Minutes': '*/5 * * * *',
    'Every 15 Minutes': '*/15 * * * *',
    'Every Hour': '0 * * * *',
    'Every Day at 9 AM': '0 9 * * *',
    'Every Weekday at 9 AM': '0 9 * * 1-5',
    'Every Monday at 9 AM': '0 9 * * 1',
};

export function validateCronExpression(cron: string): boolean {
    const parts = cron.split(' ');
    if (parts.length !== 5) return false;

    // Basic validation (could be more sophisticated)
    return parts.every(part => part.length > 0);
}

```

#### Job Persistence and Recovery:

```

# Current: In-memory scheduler (jobs lost on restart)
# Future: SQLAlchemy job store for persistence

from apscheduler.jobstores.sqlalchemy import SQLAlchemyJobStore

jobstores = {
    'default': SQLAlchemyJobStore(url=settings.DATABASE_URL)
}

scheduler = AsyncIOScheduler(
    jobstores=jobstores,

```

```

        timezone="UTC"
    )

```

**Concurrent Execution Safety:**

```

# Budget checks use row-level locking
async def check_and_reserve_budget(self, system_id: UUID, estimated_cost: float):
    # SELECT FOR UPDATE locks the row
    result = await self.db.execute(
        select(PromptSystem)
        .where(PromptSystem.id == system_id)
        .with_for_update() # Exclusive lock
    )
    system = result.scalar_one_or_none()

    # Check and update atomically
    if system.daily_cost_usd + estimated_cost > system.max_daily_cost_usd:
        raise BudgetExceededException(...)

    system.daily_cost_usd += estimated_cost
    await self.db.commit() # Lock released

```

**Resource Management:**

```

# Limit concurrent test executions
from asyncio import Semaphore

class TestExecutor:
    def __init__(self, db: AsyncSession, max_concurrent: int = 10):
        self.db = db
        self.semaphore = Semaphore(max_concurrent) # Limit concurrent API calls

    async def _execute_single_test(self, ...):
        async with self.semaphore: # Acquire semaphore
            return await self.llm_service.execute_prompt(...)
        # Semaphore released automatically

```

**Analytics & Reporting Performance Trend Analysis:**

```

// frontend/src/hooks/usePerformanceTrends.ts
export function usePerformanceTrends(systemId: string, days: number = 30) {
    return useQuery({
        queryKey: ['performance-trends', systemId, days],
        queryFn: async () => {
            const testRuns = await testRunsApi.getAll(systemId, 100);

            // Filter by date range
            const cutoff = new Date();

```

```

        cutoff.setDate(cutoff.getDate() - days);

    const recentRuns = testRuns.filter(
        run => new Date(run.started_at) >= cutoff
    );

    // Transform for chart
    return recentRuns.map(run => ({
        date: new Date(run.started_at),
        score: run.avg_score,
        cost: run.actual_cost_usd,
        model: run.model_used
    }));
}
);
}

```

### Model Comparison Algorithms:

```

// frontend/src/lib/models/analytics.ts
export function compareModels(testRuns: TestRun[]): ModelComparison {
    // Group by model
    const byModel = groupBy(testRuns, 'model_used');

    const comparisons = Object.entries(byModel).map(([model, runs]) => {
        const scores = runs.map(r => r.avg_score).filter(Boolean);
        const costs = runs.map(r => r.actual_cost_usd).filter(Boolean);
        const latencies = runs.flatMap(r =>
            r.test_results?.map(tr => tr.execution_time_ms) || []
        );

        return {
            model,
            avgScore: average(scores),
            avgCost: average(costs),
            avgLatency: average(latencies),
            totalRuns: runs.length,
            successRate: runs.filter(r => r.status === 'completed').length / runs.length
        };
    });

    return {
        comparisons,
        mostAccurate: maxBy(comparisons, 'avgScore'),
        fastest: minBy(comparisons, 'avgLatency'),
        bestValue: maxBy(comparisons, c => c.avgScore / c.avgCost)
    };
}

```

```
}
```

### Cost Optimization Insights:

```
# Analyze cost trends and suggest optimizations
def analyze_cost_trends(system_id: UUID, days: int = 30) -> CostAnalysis:
    """Analyze cost patterns and suggest optimizations."""
    test_runs = get_test_runs(system_id, days=days)

    # Group by model
    by_model = group_by(test_runs, 'model_used')

    insights = []
    for model, runs in by_model.items():
        total_cost = sum(r.actual_cost_usd for r in runs)
        avg_score = mean(r.avg_score for r in runs if r.avg_score)

        # Calculate cost per point of score
        cost_per_point = total_cost / (avg_score * len(runs)) if avg_score else float('inf')

        insights.append({
            'model': model,
            'total_cost': total_cost,
            'avg_score': avg_score,
            'cost_per_point': cost_per_point,
            'recommendation': 'optimal' if cost_per_point < threshold else 'consider_cheaper'
        })

    return CostAnalysis(insights=insights)
```

### Predictive Alerting:

```
def predict_performance_degradation(system_id: UUID) -> Optional[Alert]:
    """Predict performance issues before they occur."""
    recent_runs = get_recent_test_runs(system_id, days=7)

    if len(recent_runs) < 3:
        return None # Not enough data

    # Calculate trend (linear regression)
    scores = [r.avg_score for r in recent_runs if r.avg_score]
    if len(scores) < 3:
        return None

    # Simple linear trend
    x = list(range(len(scores)))
    slope = calculate_slope(x, scores)
```

```

# Predict next score
predicted_score = scores[-1] + slope

# Alert if predicted to drop below threshold
if predicted_score < 0.7: # Threshold
    return Alert(
        type='predicted_degradation',
        message=f"Performance predicted to drop to {predicted_score:.2f}",
        severity='warning'
    )

return None

```

---

## Technical Deep Dives

### Database Design

#### Schema Design Rationale:

##### 1. JSONB for Flexible Data:

```

-- Why JSONB instead of separate columns?
-- Flexibility: Different prompt systems have different variables
-- Extensibility: Easy to add new fields without migrations
-- Performance: PostgreSQL JSONB is indexed and queryable

CREATE TABLE prompt_systems (
    variables JSONB NOT NULL DEFAULT '[]', -- ["user_query", "context"]
    model_params JSONB NOT NULL DEFAULT '{}' -- {"temperature": 0.7, "max_tokens": 150}
);

-- Query JSONB fields
SELECT * FROM prompt_systems
WHERE model_params->>'temperature' = '0.7';

-- Index JSONB for performance
CREATE INDEX idx_model_params ON prompt_systems USING GIN (model_params);

```

##### 2. Timestamps with Timezone:

```

-- Always store timestamps with timezone (UTC)
created_at TIMESTAMP WITH TIME ZONE DEFAULT NOW()

-- Benefits:
-- - Consistent across timezones
-- - Easy to convert to local time in application
-- - Avoids daylight saving time issues

```

### 3. Decimal for Scores:

```
-- Use DECIMAL(4,3) for scores (0.000 to 1.000)
score DECIMAL(4,3) NOT NULL

-- Why not FLOAT?
-- - Exact precision (no floating point errors)
-- - Consistent storage size
-- - Better for financial-like calculations
```

### 4. Cascade Deletes:

```
-- When prompt system deleted, delete related data
CREATE TABLE regression_sets (
    system_id UUID NOT NULL REFERENCES prompt_systems(id) ON DELETE CASCADE
);

-- Prevents orphaned records
-- Ensures data consistency
```

#### Query Patterns:

##### 1. Common Queries with Indexes:

```
# Get recent test runs for a system (indexed on system_id)
result = await db.execute(
    select(TestRun)
    .where(TestRun.system_id == system_id)
    .order_by(TestRun.started_at.desc())
    .limit(50)
)
# Uses: idx_test_runs_system_id
```

##### 2. N+1 Problem Solutions:

```
# Bad: N+1 queries
test_runs = await db.execute(select(TestRun))
for run in test_runs:
    results = await db.execute(
        select(TestResult).where(TestResult.test_run_id == run.id)
    ) # One query per run!

# Good: Eager loading
from sqlalchemy.orm import selectinload

test_runs = await db.execute(
    select(TestRun)
    .options(selectinload(TestRun.test_results)) # Load in one query
)
```

### 3. Pagination Strategies:

```
# Offset-based pagination (simple but slow for large offsets)
def get_test_runs_page(offset: int = 0, limit: int = 50):
    return select(TestRun).offset(offset).limit(limit)

# Cursor-based pagination (faster, better for large datasets)
def get_test_runs_after(cursor: datetime, limit: int = 50):
    return select(TestRun).where(
        TestRun.started_at > cursor
    ).order_by(TestRun.started_at).limit(limit)
```

### 4. Caching Implementation (Future):

```
# Redis caching for frequently accessed data
from redis import asyncio as aioredis

redis = await aioredis.from_url("redis://localhost")

async def get_prompt_system_cached(system_id: UUID):
    # Check cache first
    cached = await redis.get(f"system:{system_id}")
    if cached:
        return json.loads(cached)

    # Load from database
    system = await db.execute(...)

    # Cache for 5 minutes
    await redis.setex(
        f"system:{system_id}",
        300, # 5 minutes
        json.dumps(system)
    )

    return system
```

## API Design

### RESTful Endpoints:

```
# Prompt Systems
GET /api/prompt-systems           # List all
POST /api/prompt-systems          # Create
GET /api/prompt-systems/{id}       # Get one
PUT /api/prompt-systems/{id}       # Update
DELETE /api/prompt-systems/{id}    # Delete
```

```

# Regression Sets
GET    /api/prompt-systems/{id}/regression-sets      # List
POST   /api/prompt-systems/{id}/regression-sets      # Create
DELETE /api/prompt-systems/{id}/regression-sets/{reg_id} # Delete

# Test Runs
POST   /api/test-runs                               # Create (returns 202 Accepted)
GET    /api/test-runs                                # List (with filters)
GET    /api/test-runs/{id}                            # Get status
GET    /api/test-runs/{id}/results                 # Get detailed results

# Schedules
GET    /api/schedules                                # List
POST   /api/schedules                                # Create
PUT    /api/schedules/{id}                            # Update
DELETE /api/schedules/{id}                            # Delete
POST   /api/schedules/{id}/trigger                 # Manual trigger

```

#### Request/Response Examples:

```

// POST /api/prompt-systems
Request:
{
  "name": "Customer Support Assistant",
  "description": "Handles customer inquiries",
  "template": "You are a support agent. Question: {user_query}",
  "variables": ["user_query"],
  "model": "gpt-3.5-turbo",
  "model_params": {
    "temperature": 0.7,
    "max_tokens": 150
  }
}

Response (201 Created):
{
  "id": "550e8400-e29b-41d4-a716-446655440000",
  "name": "Customer Support Assistant",
  "template": "You are a support agent. Question: {user_query}",
  "variables": ["user_query"],
  "model": "gpt-3.5-turbo",
  "model_params": {
    "temperature": 0.7,
    "max_tokens": 150
  },
  "created_at": "2024-10-15T10:30:00Z"
}

```

```

// POST /api/test-runs
Request:
{
  "system_id": "550e8400-e29b-41d4-a716-446655440000",
  "model_override": "gpt-4o",
  "triggered_by": "manual"
}

Response (202 Accepted):
{
  "message": "Test run started",
  "test_run_id": "660e8400-e29b-41d4-a716-446655440001",
  "status": "pending"
}

// GET /api/test-runs/{id}
Response (200 OK):
{
  "id": "660e8400-e29b-41d4-a716-446655440001",
  "system_id": "550e8400-e29b-41d4-a716-446655440000",
  "status": "running",
  "model_used": "gpt-4o",
  "total_tests": 3,
  "completed_tests": 2,
  "avg_score": null, // Not calculated yet
  "started_at": "2024-10-15T10:35:00Z",
  "completed_at": null
}

```

#### Error Codes and Messages:

```

# Standard HTTP status codes
400 Bad Request: Validation error, invalid input
401 Unauthorized: Missing or invalid authentication (future)
402 Payment Required: Budget exceeded
404 Not Found: Resource not found
422 Unprocessable: Business logic error
500 Internal Server: Unexpected server error

```

```

# Error response format
{
  "detail": "Human-readable error message",
  "code": "BUDGET_EXCEEDED", # Machine-readable code
  "field": "system_id" # Optional: field that caused error
}

```

### Rate Limiting (Future):

```
from slowapi import Limiter
from slowapi.util import get_remote_address

limiter = Limiter(key_func=get_remote_address)

@app.post("/api/test-runs")
@limiter.limit("10/minute") # 10 requests per minute per IP
async def create_test_run(...):
    ...
```

### Security Implementation

#### API Key Management:

```
# Environment variables (never in code)
# backend/.env
OPENAI_API_KEY=sk-proj-xxxxxxxxxxxxxxxxxxxxxx

# Loaded via Pydantic settings
class Settings(BaseSettings):
    OPENAI_API_KEY: str
    class Config:
        env_file = ".env"

# Never log API keys
logger.info(f"Calling OpenAI API") # OK
logger.info(f"API key: {settings.OPENAI_API_KEY}") # NEVER
```

#### Input Validation and Sanitization:

```
# Pydantic automatically validates and sanitizes
class PromptSystemCreate(BaseModel):
    name: str = Field(..., min_length=1, max_length=255)
    template: str = Field(..., min_length=1)

    @validator('name')
    def validate_name(cls, v):
        # Remove leading/trailing whitespace
        v = v.strip()
        # Prevent SQL injection (Pydantic handles this)
        if not v:
            raise ValueError("Name cannot be empty")
        return v
```

#### Audit Logging:

```
# Log all test executions for audit trail
```

```

async def execute_test_run(...):
    logger.info(
        f"Test run started: system_id={system_id}, "
        f"model={model}, triggered_by={triggered_by}, "
        f"user_ip={request.client.host}" # If auth added
    )

    # Store audit log in database
    audit_log = AuditLog(
        action="test_run_created",
        system_id=system_id,
        user_id=user_id, # Future: when auth added
        metadata={"model": model, "triggered_by": triggered_by}
    )
    db.add(audit_log)

```

---

## Performance & Scalability

### Performance Characteristics

#### Benchmark Results:

**Test Execution Times:** - Single test (GPT-3.5-turbo): 2-3 seconds - Single test (GPT-4o): 2-3 seconds  
- Single test (GPT-4o-mini): 1-2 seconds - 10-test regression set (GPT-3.5-turbo): 15-20 seconds (sequential) - 10-test regression set (parallel): 3-5 seconds (theoretical)

**API Response Times:** - GET /api/prompt-systems: < 50ms (10 systems)  
- GET /api/test-runs: < 100ms (50 runs) - POST /api/test-runs: < 200ms (returns immediately, execution in background) - Database queries: < 50ms (indexed queries)

**UI Rendering:** - Initial page load: < 500ms - Component render: < 50ms - Chart rendering: < 100ms (Recharts)

**Memory Usage:** - Backend: ~100MB base + ~10MB per concurrent test run  
- Frontend: ~5MB bundle size, ~50MB runtime memory

#### Optimization Strategies:

##### 1. Database Indexing:

```

-- Critical indexes for performance
CREATE INDEX idx_test_runs_system_id ON test_runs(system_id);
CREATE INDEX idx_test_runs_started_at ON test_runs(started_at DESC);
CREATE INDEX idx_test_results_test_run_id ON test_results(test_run_id);
CREATE INDEX idx_prompt_systems_last_reset_date ON prompt_systems(last_reset_date);

```

```
-- Composite index for common query pattern
CREATE INDEX idx_test_runs_system_status ON test_runs(system_id, status);
```

## 2. Connection Pooling:

```
# Optimize connection pool for expected load
engine = create_async_engine(
    settings.DATABASE_URL,
    pool_size=20,           # Base connections
    max_overflow=40,         # Additional under load
    pool_pre_ping=True,     # Verify connections
    pool_recycle=3600        # Recycle after 1 hour
)
```

## 3. Async/Await Patterns:

```
# Sequential (current)
for regression_set in regression_sets:
    result = await execute_test(regression_set) # Waits for each

# Parallel (future optimization)
results = await asyncio.gather(*[
    execute_test(rs) for rs in regression_sets
]) # Executes all concurrently
```

## 4. Caching (Future):

```
# Cache frequently accessed data
@lru_cache(maxsize=100)
def get_model_pricing(model: str):
    return MODEL_PRICING[model]

# Redis for shared cache
async def get_prompt_system_cached(system_id: UUID):
    cache_key = f"system:{system_id}"
    cached = await redis.get(cache_key)
    if cached:
        return json.loads(cached)
    # ... load from DB and cache
```

## Scalability Considerations

### Horizontal Scaling Architecture:

Load Balancer  
(Nginx/HAProxy)

```
FastAPI      FastAPI      FastAPI  
Server       Server       Server  
#1          #2          #3
```

PostgreSQL  
(Primary)

- Stateless Backend:**
- No session state stored in memory - All state in database
  - Any server can handle any request - Easy horizontal scaling

#### Database Sharding (Future):

```
# Shard by system_id (hash-based)
def get_shard(system_id: UUID) -> str:
    hash_value = hash(str(system_id))
    shard_num = hash_value % NUM_SHARDS
    return f"shard_{shard_num}"

# Route queries to appropriate shard
async def get_test_runs(system_id: UUID):
    shard = get_shard(system_id)
    db = get_shard_connection(shard)
    return await db.execute(...)
```

#### Microservices Migration Path:

```
Current: Monolithic FastAPI app
↓
Phase 1: Extract scheduler to separate service
- Scheduler Service (APScheduler)
- Main API Service (FastAPI)
- Shared database
↓
Phase 2: Extract test execution to worker pool
- API Service (receives requests)
- Worker Service (executes tests, Celery/Redis)
- Scheduler Service (triggers workers)
↓
Phase 3: Full microservices
```

- Prompt System Service
- Test Execution Service
- Analytics Service
- Notification Service

#### Load Balancing Strategies:

```
# Nginx configuration for load balancing
upstream fastapi_backend {
    least_conn; # Use least connections algorithm
    server api1:8000;
    server api2:8000;
    server api3:8000;
}

server {
    location /api/ {
        proxy_pass http://fastapi_backend;
        proxy_set_header Host $host;
        proxy_set_header X-Real-IP $remote_addr;
    }
}
```

#### Database Replication:

```
-- Read replicas for scaling reads
-- Primary: Handles writes
-- Replicas: Handle reads (test run queries, analytics)

-- Application code routes reads to replicas
async def get_test_runs(system_id: UUID):
    db = get_read_replica() # Round-robin or least-lag
    return await db.execute(...)

async def create_test_run(...):
    db = get_primary() # Always use primary for writes
    ...
```

---

#### End of Part 2

*Continue to Part 3 for: Testing & QA, Deployment & Operations, Challenges & Solutions, Future Roadmap, Appendices*

# AI Prompt System Monitor - Technical Specification

## Part 3: Testing, Deployment, Challenges & Future Roadmap

**Document Version:** 1.0

**Date:** 2024

**Author:** Yi Zu

**Project:** AI Prompt System Monitor - Production-Ready Full-Stack Application

---

### Table of Contents

1. Testing & Quality Assurance
  2. Deployment & Operations
  3. Challenges & Solutions
  4. Future Roadmap
  5. Appendices
- 

## Testing & Quality Assurance

### Testing Strategy

#### Unit Testing Service Layer Testing Examples:

```
# tests/test_llm_service.py
import pytest
from app.services.llm_service import LLMService

class TestLLMService:
    def test_interpolate_template(self):
        """Test template variable interpolation."""
        template = "Hello {name}, you have {count} messages."
        variables = {"name": "Alice", "count": 5}

        result = LLMService.interpolate_template(template, variables)

        assert result == "Hello Alice, you have 5 messages."

    def test_interpolate_template_missing_variable(self):
        """Test that missing variables are handled."""
        template = "Hello {name}"
        variables = {} # Missing 'name'
```

```

        result = LLMService.interpolate_template(template, variables)

    assert result == "Hello {name}"  # Unchanged if variable missing

def test_calculate_similarity_exact_match(self):
    """Test similarity calculation for exact match."""
    actual = "We are open Monday-Friday 9am-5pm PT"
    expected = "We are open Monday-Friday 9am-5pm PT"

    score = LLMService.calculate_similarity(actual, expected)

    assert score == 1.0

def test_calculate_similarity_substring_match(self):
    """Test similarity for substring match."""
    actual = "Our hours: We are open Monday-Friday 9am-5pm PT. Thanks!"
    expected = "We are open Monday-Friday 9am-5pm PT"

    score = LLMService.calculate_similarity(actual, expected)

    assert score == 0.95

def test_calculate_similarity_word_overlap(self):
    """Test similarity for word overlap."""
    actual = "We operate Monday through Friday, 9am to 5pm"
    expected = "We are open Monday-Friday 9am-5pm PT"

    score = LLMService.calculate_similarity(actual, expected)

    assert 0.0 < score < 1.0  # Some overlap but not perfect
    assert score > 0.1  # Should have meaningful overlap

```

### Mock Implementations:

```

# tests/conftest.py
import pytest
from unittest.mock import AsyncMock, MagicMock
from app.services.llm_service import LLMService

@pytest.fixture
def mock_openai_client():
    """Mock OpenAI client for testing."""
    client = MagicMock()
    client.chat.completions.create = AsyncMock(return_value=MagicMock(
        choices=[MagicMock(message=MagicMock(content="Mocked response"))],
        usage=MagicMock()
    ))

```

```

        prompt_tokens=100,
        completion_tokens=50,
        total_tokens=150
    )
))
return client

@pytest.fixture
def llm_service(mock_openai_client):
    """LLMService with mocked OpenAI client."""
    service = LLMService()
    service.client = mock_openai_client
    return service

# Usage in tests
async def test_execute_prompt(llm_service, mock_openai_client):
    """Test prompt execution with mocked OpenAI."""
    result = await llm_service.execute_prompt(
        template="Test {variable}",
        variables={"variable": "value"},
        model="gpt-4o"
    )

    assert result.output == "Mocked response"
    assert result.total_tokens == 150
    mock_openai_client.chat.completions.create.assert_called_once()

```

#### Test Coverage Metrics:

```

# Run tests with coverage
pytest --cov=app --cov-report=html --cov-report=term

# Target coverage:
# - Services: 80%+ (business logic critical)
# - API routes: 70%+ (happy paths + error cases)
# - Models: 60%+ (mostly ORM, less critical)

```

#### Integration Testing API Endpoint Testing:

```

# tests/test_api.py
import pytest
from fastapi.testclient import TestClient
from app.main import app

client = TestClient(app)

def test_create_prompt_system():

```

```

"""Test creating a prompt system via API."""
response = client.post(
    "/api/prompt-systems",
    json={
        "name": "Test System",
        "template": "Hello {name}",
        "variables": ["name"],
        "model": "gpt-3.5-turbo"
    }
)

assert response.status_code == 201
data = response.json()
assert data["name"] == "Test System"
assert "id" in data

def test_create_test_run():
    """Test creating a test run."""
    # First create a system and regression set
    system_response = client.post("/api/prompt-systems", json={...})
    system_id = system_response.json()["id"]

    # Create regression set
    client.post(f"/api/prompt-systems/{system_id}/regression-sets", json={...})

    # Create test run
    response = client.post(
        "/api/test-runs",
        json={"system_id": system_id}
    )

    assert response.status_code == 202 # Accepted
    assert "test_run_id" in response.json()

```

#### Database Integration Tests:

```

# tests/test_database.py
import pytest
from sqlalchemy.ext.asyncio import AsyncSession
from app.core.database import AsyncSessionLocal
from app.models.prompt_system import PromptSystem

@pytest.mark.asyncio
async def test_create_prompt_system_in_db():
    """Test creating prompt system in database."""
    async with AsyncSessionLocal() as db:
        system = PromptSystem(

```

```

        name="Test System",
        template="Hello {name}",
        variables=["name"],
        model="gpt-3.5-turbo"
    )
db.add(system)
await db.commit()
await db.refresh(system)

assert system.id is not None
assert system.name == "Test System"

```

#### External Service Mocking:

```

# Mock OpenAI API responses
@pytest.fixture
def mock_openai_response():
    """Mock OpenAI API response."""
    return {
        "choices": [
            {
                "message": {
                    "content": "Mocked AI response"
                }
            }
        ],
        "usage": {
            "prompt_tokens": 100,
            "completion_tokens": 50,
            "total_tokens": 150
        }
    }

# Use httpx to mock HTTP requests
from httpx import AsyncClient
import respx

@respx.mock
async def test_openai_integration():
    """Test OpenAI integration with mocked HTTP."""
    respx.post("https://api.openai.com/v1/chat/completions").mock(
        return_value=httpx.Response(200, json=mock_openai_response())
    )

    service = LLMService()
    result = await service.execute_prompt("Test", {}, "gpt-4o")

    assert result.output == "Mocked AI response"

```

### End-to-End Testing User Journey Testing:

```
# tests/e2e/test_user_workflow.py
def test_complete_user_workflow():
    """Test complete workflow: create system + add tests + run + view results."""
    # 1. Create prompt system
    system_response = client.post("/api/prompt-systems", json={...})
    system_id = system_response.json()["id"]

    # 2. Add regression test
    client.post(f"/api/prompt-systems/{system_id}/regression-sets", json={...})

    # 3. Run tests
    run_response = client.post("/api/test-runs", json={"system_id": system_id})
    test_run_id = run_response.json()["test_run_id"]

    # 4. Poll for completion
    import time
    for _ in range(30): # Max 30 seconds
        status_response = client.get(f"/api/test-runs/{test_run_id}")
        status = status_response.json()["status"]
        if status == "completed":
            break
        time.sleep(1)

    assert status == "completed"

    # 5. Get results
    results_response = client.get(f"/api/test-runs/{test_run_id}/results")
    results = results_response.json()

    assert len(results) > 0
    assert results[0]["score"] is not None
```

### Performance Testing:

```
# tests/performance/test_load.py
import asyncio
import time

async def test_concurrent_test_runs():
    """Test system under concurrent load."""
    # Create 10 test runs simultaneously
    tasks = [
        create_test_run(system_id)
        for _ in range(10)
    ]
```

```

start_time = time.time()
results = await asyncio.gather(*tasks)
duration = time.time() - start_time

# All should complete successfully
assert all(r["status"] == "pending" for r in results)

# Should complete quickly (just creating records)
assert duration < 2.0 # 2 seconds for 10 concurrent requests

```

### Chaos Engineering:

```

# tests/chaos/test_failure_modes.py
def test_openai_api_failure():
    """Test behavior when OpenAI API fails."""
    # Simulate API failure
    with patch('app.services.llm_service.AsyncOpenAI') as mock:
        mock.return_value.chat.completions.create.side_effect = Exception("API Error")

        # Test run should handle gracefully
        response = client.post("/api/test-runs", json={...})
        test_run_id = response.json()["test_run_id"]

        # Wait for failure
        time.sleep(5)

        status_response = client.get(f"/api/test-runs/{test_run_id}")
        assert status_response.json()["status"] == "failed"

def test_database_connection_loss():
    """Test behavior when database connection is lost."""
    # Simulate DB failure
    # System should return 503 Service Unavailable
    # Not crash or hang

```

### Quality Metrics

#### Code Quality Standards:

```

# .flake8 or pyproject.toml
[tool.flake8]
max-line-length = 100
exclude = ["venv", "migrations"]
ignore = ["E203", "W503"] # Allow some formatting

[tool.black]

```

```

line-length = 100
target-version = ['py310']

[tool.mypy]
python_version = "3.10"
warn_return_any = True
warn_unused_configs = True

Performance Benchmarks:

# tests/benchmarks/test_performance.py
def benchmark_similarity_calculation(benchmark):
    """Benchmark similarity calculation speed."""
    actual = "We are open Monday-Friday 9am-5pm PT"
    expected = "We operate Monday through Friday, 9am to 5pm"

    result = benchmark(LLMService.calculate_similarity, actual, expected)

    # Should complete in < 1ms
    assert result < 0.001

def benchmark_database_query(benchmark):
    """Benchmark database query performance."""
    async def query():
        async with AsyncSessionLocal() as db:
            result = await db.execute(select(TestRun).limit(50))
            return result.scalars().all()

    result = benchmark(query)

    # Should complete in < 100ms
    assert result < 0.1

Security Assessments:

# tests/security/test_security.py
def test_sql_injection_prevention():
    """Test that SQL injection is prevented."""
    # Attempt SQL injection in name field
    malicious_input = "'; DROP TABLE prompt_systems; --"

    response = client.post(
        "/api/prompt-systems",
        json={"name": malicious_input, "template": "Test", ...}
    )

    # Should either reject (validation) or escape (safe)
    # Database should still exist

```

```

assert response.status_code in [400, 201] # Rejected or safely handled

def test_xss_prevention():
    """Test that XSS is prevented in output."""
    # Store script tag in template
    response = client.post(
        "/api/prompt-systems",
        json={"template": "<script>alert('xss')</script>", ...}
    )

    # Frontend should escape when rendering
    # Or backend should sanitize

```

#### Maintainability Metrics:

- **Cyclomatic Complexity:** < 10 per function
  - **Function Length:** < 50 lines
  - **Class Length:** < 300 lines
  - **Test Coverage:** > 70% overall
  - **Documentation:** All public functions documented
- 

## Deployment & Operations

### Infrastructure Requirements

#### Development Environment Local Development Setup:

```

# Prerequisites
- Python 3.10+
- Node.js 18+
- PostgreSQL 12+ (or Supabase account)
- OpenAI API key

# Backend setup
cd backend
python3 -m venv venv
source venv/bin/activate
pip install -r requirements.txt

# Frontend setup
cd frontend
npm install

# Database setup
# Option 1: Local PostgreSQL
createdb ai_prompt_monitor

```

```

psql ai_prompt_monitor < DATABASE_SCHEMA.sql

# Option 2: Supabase (recommended for demo)
# - Create project at supabase.com
# - Run DATABASE_SCHEMA.sql in SQL Editor
# - Copy connection string to backend/.env

```

#### Database Provisioning:

```

-- Development database setup
CREATE DATABASE ai_prompt_monitor;
\c ai_prompt_monitor

-- Run schema
\i DATABASE_SCHEMA.sql

-- Create test user (optional)
CREATE USER dev_user WITH PASSWORD 'dev_password';
GRANT ALL PRIVILEGES ON DATABASE ai_prompt_monitor TO dev_user;

```

#### API Key Management:

```

# backend/.env (never commit to git)
DATABASE_URL=postgresql+asyncpg://user:pass@localhost:5432/ai_prompt_monitor
OPENAI_API_KEY=sk-proj-xxxxxxxxxxxxxxxxxxxxxx
BACKEND_HOST=0.0.0.0
BACKEND_PORT=8000
FRONTEND_URL=http://localhost:5173
ALLOWED_ORIGINS=http://localhost:5173

# Add to .gitignore
echo ".env" >> .gitignore

```

#### Production Environment Cloud Provider Recommendations:

**Option 1: Vercel (Frontend) + Railway (Backend) + Supabase (Database)** - **Frontend:** Vercel (free tier, automatic deployments) - **Backend:** Railway (easy PostgreSQL, \$5/month) - **Database:** Supabase (free tier, managed PostgreSQL) - **Pros:** Easy setup, good free tiers - **Cons:** Vendor lock-in, scaling costs

**Option 2: AWS (Full Stack)** - **Frontend:** S3 + CloudFront (static hosting) - **Backend:** ECS Fargate or Lambda (serverless) - **Database:** RDS PostgreSQL - **Pros:** Scalable, enterprise-grade - **Cons:** Complex setup, higher cost

**Option 3: Docker + Kubernetes** - **Frontend:** Containerized React app - **Backend:** Containerized FastAPI - **Database:** Managed PostgreSQL (Cloud SQL, RDS) - **Orchestration:** Kubernetes (GKE, EKS, AKS) - **Pros:** Portable, scalable - **Cons:** Requires Kubernetes expertise

### Container Orchestration:

```
# Dockerfile (Backend)
FROM python:3.10-slim

WORKDIR /app

# Install dependencies
COPY requirements.txt .
RUN pip install --no-cache-dir -r requirements.txt

# Copy application
COPY .

# Run application
CMD ["uvicorn", "app.main:app", "--host", "0.0.0.0", "--port", "8000"]

# Dockerfile (Frontend)
FROM node:18-alpine AS builder

WORKDIR /app
COPY package*.json .
RUN npm ci
COPY ..
RUN npm run build

FROM nginx:alpine
COPY --from=builder /app/dist /usr/share/nginx/html
COPY nginx.conf /etc/nginx/nginx.conf

EXPOSE 80
CMD ["nginx", "-g", "daemon off;"]

# docker-compose.yml (Local development)
version: '3.8'

services:
  backend:
    build: ./backend
    ports:
      - "8000:8000"
    environment:
      - DATABASE_URL=postgresql+asyncpg://postgres:password@db:5432/ai_prompt_monitor
      - OPENAI_API_KEY=${OPENAI_API_KEY}
    depends_on:
      - db
```

```

frontend:
  build: ./frontend
  ports:
    - "80:80"
depends_on:
  - backend

db:
  image: postgres:14
  environment:
    - POSTGRES_DB=ai_prompt_monitor
    - POSTGRES_PASSWORD=password
  volumes:
    - postgres_data:/var/lib/postgresql/data

volumes:
  postgres_data:

```

## CI/CD Pipeline

### Build Automation:

```

# .github/workflows/ci.yml
name: CI/CD Pipeline

on:
  push:
    branches: [main, develop]
  pull_request:
    branches: [main]

jobs:
  test-backend:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v3
      - uses: actions/setup-python@v4
        with:
          python-version: '3.10'
      - name: Install dependencies
        run: |
          cd backend
          pip install -r requirements.txt
          pip install pytest pytest-cov
      - name: Run tests
        run: |

```

```

        cd backend
        pytest --cov=app --cov-report=xml
- name: Upload coverage
  uses: codecov/codecov-action@v3

test-frontend:
  runs-on: ubuntu-latest
  steps:
    - uses: actions/checkout@v3
    - uses: actions/setup-node@v3
      with:
        node-version: '18'
    - name: Install dependencies
      run: |
        cd frontend
        npm ci
    - name: Run tests
      run: |
        cd frontend
        npm test
    - name: Build
      run: |
        cd frontend
        npm run build

deploy:
  needs: [test-backend, test-frontend]
  runs-on: ubuntu-latest
  if: github.ref == 'refs/heads/main'
  steps:
    - uses: actions/checkout@v3
    - name: Deploy to production
      run: |
        # Deploy backend to Railway
        # Deploy frontend to Vercel
        # Run database migrations

```

#### Testing Automation:

```

# Automated testing in CI
- name: Run integration tests
  run: |
    # Start test database
    docker-compose up -d db
    # Wait for DB
    sleep 5
    # Run tests

```

```
pytest tests/integration/
```

#### Deployment Strategies:

**Blue-Green Deployment:** 1. Deploy new version to “green” environment 2. Run smoke tests 3. Switch traffic from “blue” to “green” 4. Monitor for errors 5. Rollback to “blue” if issues detected

#### Rollback Procedures:

```
# Rollback script
#!/bin/bash

# 1. Identify current version
CURRENT_VERSION=$(git describe --tags)

# 2. Revert to previous version
git checkout previous-tag

# 3. Rebuild and redeploy
docker-compose up -d --build

# 4. Verify health
curl http://localhost:8000/health

# 5. Notify team
echo "Rolled back to $CURRENT_VERSION"
```

### Monitoring and Logging

#### Application Monitoring:

```
# backend/app/main.py
import logging
from prometheus_client import Counter, Histogram, generate_latest

# Metrics
test_runs_total = Counter('test_runs_total', 'Total test runs')
test_execution_time = Histogram('test_execution_seconds', 'Test execution time')
api_requests = Counter('api_requests_total', 'API requests', ['method', 'endpoint'])

@app.middleware("http")
async def metrics_middleware(request: Request, call_next):
    start_time = time.time()
    response = await call_next(request)
    duration = time.time() - start_time

    api_requests.labels(method=request.method, endpoint=request.url.path).inc()
```

```

        return response

@app.get("/metrics")
async def metrics():
    """Prometheus metrics endpoint."""
    return Response(generate_latest(), media_type="text/plain")

```

#### Logging Configuration:

```

# backend/app/core/logging.py
import logging
import sys
from logging.handlers import RotatingFileHandler

def setup_logging():
    """Configure application logging."""
    logging.basicConfig(
        level=logging.INFO,
        format='%(asctime)s - %(name)s - %(levelname)s - %(message)s',
        handlers=[
            RotatingFileHandler('app.log', maxBytes=10*1024*1024, backupCount=5),
            logging.StreamHandler(sys.stdout)
        ]
    )

    # Set log levels for libraries
    logging.getLogger('sqlalchemy.engine').setLevel(logging.WARNING)
    logging.getLogger('httpx').setLevel(logging.WARNING)

```

#### Health Checks:

```

# backend/app/main.py
@app.get("/health")
async def health():
    """Comprehensive health check."""
    checks = {
        "status": "healthy",
        "database": "unknown",
        "scheduler": "unknown",
        "openai": "unknown"
    }

    # Check database
    try:
        async with AsyncSessionLocal() as db:
            await db.execute(text("SELECT 1"))
        checks["database"] = "connected"
    
```

```

except Exception as e:
    checks["database"] = f"error: {str(e)}"
    checks["status"] = "unhealthy"

# Check scheduler
if scheduler.scheduler.running:
    checks["scheduler"] = "running"
else:
    checks["scheduler"] = "stopped"
    checks["status"] = "unhealthy"

# Check OpenAI (lightweight check)
try:
    # Could ping OpenAI API or check API key format
    if settings.OPENAI_API_KEY.startswith("sk-"):
        checks["openai"] = "configured"
except:
    checks["openai"] = "not_configured"

status_code = 200 if checks["status"] == "healthy" else 503
return JSONResponse(content=checks, status_code=status_code)

```

---

## Challenges & Solutions

### Technical Challenges

Async Execution Complexity   Challenge: Race Conditions in Budget Checks

**Problem:** - Multiple scheduled tests run simultaneously - Both check budget before either updates - Both pass budget check, exceeding limit

**Solution: Row-Level Locking**

```

# backend/app/services/budget_guard.py
async def check_and_reserve_budget(self, system_id: UUID, estimated_cost: float):
    # SELECT FOR UPDATE locks the row exclusively
    result = await self.db.execute(
        select(PromptSystem)
        .where(PromptSystem.id == system_id)
        .with_for_update()  # Exclusive lock
    )
    system = result.scalar_one_or_none()

    # Check and update atomically
    if system.daily_cost_usd + estimated_cost > system.max_daily_cost_usd:

```

```

    raise BudgetExceededException(...)

    system.daily_cost_usd += estimated_cost
    await self.db.commit()  # Lock released

    # Other transactions wait here until lock released

```

### Challenge: Resource Cleanup

**Problem:** - Background tasks may fail - Database sessions not properly closed  
- Memory leaks from unclosed connections

### Solution: Context Managers

```

# Always use context managers for database sessions
async def _execute_test_run_background(test_run_id: UUID):
    async with AsyncSessionLocal() as db:  # Auto-closes on exit
        try:
            executor = TestExecutor(db)
            await executor.execute_test_run_by_id(test_run_id)
        except Exception as e:
            logger.error(f"Execution failed: {e}")
            # Session automatically rolled back and closed

```

### Challenge: Error Propagation

**Problem:** - Errors in background tasks not visible to API caller - Test runs stuck in “running” status

### Solution: Status Updates and Logging

```

async def execute_test_run_by_id(self, test_run_id: UUID):
    test_run = await self._load_test_run(test_run_id)
    test_run.status = "running"
    await self.db.commit()

    try:
        # Execute tests...
        test_run.status = "completed"
    except Exception as e:
        logger.error(f"Test run failed: {e}", exc_info=True)
        test_run.status = "failed"
        test_run.error_message = str(e)
    finally:
        await self.db.commit()

```

### Cost Management Challenge: Budget Enforcement Accuracy

**Problem:** - Estimated costs may differ from actual costs - Budget could be exceeded if estimation is too low

### Solution: Safety Buffer + Adjustment

```
# Estimate with 20% buffer
estimated_cost = CostCalculator.estimate_cost(model, num_tests) * 1.2

# Reserve estimated cost
await budget_guard.check_and_reserve_budget(system_id, estimated_cost)

# After execution, adjust for actual cost
actual_cost = sum(result.cost_usd for result in results)
cost_difference = actual_cost - estimated_cost

# Adjust daily cost
system.daily_cost_usd += cost_difference # Could be negative (refund)
```

### Challenge: Usage Monitoring

**Problem:** - Need real-time visibility into costs - Historical cost analysis

### Solution: Comprehensive Cost Tracking

```
# Store costs at multiple levels
- test_results.cost_usd (per test)
- test_runs.actual_cost_usd (per run)
- prompt_systems.daily_cost_usd (daily aggregate)
- prompt_systems.total_cost_usd (lifetime)

# Query for analytics
SELECT
    DATE(started_at) as date,
    SUM(actual_cost_usd) as daily_cost
FROM test_runs
WHERE system_id = $1
GROUP BY DATE(started_at)
ORDER BY date DESC;
```

### Real-time Updates Challenge: Polling vs WebSockets Trade-offs

**Problem:** - Polling: Simple but inefficient, slight delay - WebSockets: Real-time but complex, connection management

### Solution: Polling with Smart Intervals

```
// Adaptive polling interval
const pollTestRun = async (testRunId: string) => {
    let interval = 2000; // Start with 2 seconds

    while (true) {
        const testRun = await testRunsApi.getById(testRunId);
```

```

    if (testRun.status === 'completed' || testRun.status === 'failed') {
      return testRun;
    }

    // Adaptive interval: faster when running, slower when pending
    if (testRun.status === 'running') {
      interval = 1000; // 1 second when active
    } else {
      interval = 2000; // 2 seconds when pending
    }

    await new Promise(resolve => setTimeout(resolve, interval));
  }
};

Future: WebSocket Implementation

```

```

# backend/app/main.py
from fastapi import WebSocket

@app.websocket("/ws/test-runs/{test_run_id}")
async def websocket_test_run(websocket: WebSocket, test_run_id: UUID):
    await websocket.accept()

    try:
        while True:
            test_run = await get_test_run(test_run_id)
            await websocket.send_json({
                "status": test_run.status,
                "completed_tests": test_run.completed_tests,
                "total_tests": test_run.total_tests
            })

            if test_run.status in ["completed", "failed"]:
                break

            await asyncio.sleep(1)
    except WebSocketDisconnect:
        pass

```

### Challenge: Progress Tracking

**Problem:** - Need to show progress during long-running test execution - Multiple tests executing, need per-test progress

### Solution: Incremental Status Updates

```
# Update progress after each test
```

```

for regression_set in regression_sets:
    test_result = await self._execute_single_test(...)
    test_run.completed_tests += 1
    await self.db.commit() # Make progress visible immediately

# Frontend polls and sees: 1/10, 2/10, 3/10, ...

```

## Business Challenges

### Challenge: User Adoption

**Problem:** - Developers may not see immediate value - Learning curve for new tool

**Solution:** - Clear documentation and examples - Quick start guide (5-minute setup) - Demo video showing value - Integration with existing workflows

### Challenge: Feature Prioritization

**Problem:** - Many potential features (Level 3, multi-provider, etc.) - Limited development resources

**Solution:** - User feedback collection - Analytics on feature usage - MVP approach: Build core features first - Iterate based on real usage

### Challenge: Market Positioning

**Problem:** - Competing with commercial solutions - Need to differentiate

**Solution:** - Open-source approach (community contributions) - Focus on developer experience - Cost transparency (unlike some commercial tools) - Extensibility (custom evaluation functions)

---

## Future Roadmap

### Phase 3 Implementation (Genetic Optimization)

#### Evolutionary Algorithms for Prompt Improvement:

```

# Future: Genetic prompt optimization
class GeneticPromptOptimizer:
    def __init__(self, initial_prompt: str, budget: float):
        self.initial_prompt = initial_prompt
        self.budget = budget
        self.population = self._initialize_population(initial_prompt)

    def optimize(self) -> str:
        """Evolve prompt to maximize similarity scores."""
        for generation in range(self.max_generations):

```

```

# Evaluate fitness (run tests, calculate scores)
fitness_scores = self._evaluate_population()

# Select best performers
parents = self._select_parents(fitness_scores)

# Crossover and mutation
offspring = self._breed(parents)

# Replace population
self.population = self._replace_population(offspring, fitness_scores)

# Check budget
if self._total_cost() > self.budget:
    break

return self._best_prompt()

def _mutate_prompt(self, prompt: str) -> str:
    """Apply mutations: add words, remove words, rephrase."""
    # Random mutations to explore search space
    ...

```

### Budget-Aware Optimization:

```

def optimize_with_budget(
    self,
    prompt: str,
    max_cost: float,
    target_score: float = 0.9
) -> OptimizedPrompt:
    """Optimize prompt while respecting budget."""
    cost_tracker = CostTracker(max_cost)

    while cost_tracker.remaining() > 0:
        candidate = self._generate_candidate(prompt)
        score, cost = self._evaluate(candidate)

        if score >= target_score:
            return OptimizedPrompt(candidate, score, cost)

        cost_tracker.record(cost)

    return OptimizedPrompt(prompt, 0.0, 0.0)  # Budget exhausted

```

### A/B Testing Framework:

```
# Future: A/B testing for prompt variants
```

```

class ABTestFramework:
    async def create_ab_test(
        self,
        variant_a: str,
        variant_b: str,
        test_cases: List[RegressionSet]
    ) -> ABTestResult:
        """Run A/B test comparing two prompt variants."""
        # Run tests for both variants
        results_a = await self._run_tests(variant_a, test_cases)
        results_b = await self._run_tests(variant_b, test_cases)

        # Statistical analysis
        return ABTestResult(
            variant_a_score=mean(results_a.scores),
            variant_b_score=mean(results_b.scores),
            statistical_significance=self._calculate_significance(results_a, results_b),
            recommendation=self._recommend_variant(results_a, results_b)
        )

```

## Advanced Features

### Multi-Provider AI Support:

```

# Future: Support multiple AI providers
class MultiProviderLLMService:
    def __init__(self):
        self.providers = {
            "openai": OpenAIProvider(),
            "anthropic": AnthropicProvider(),
            "mistral": MistralProvider(),
            "cohere": CohereProvider()
        }

    async def execute_prompt(
        self,
        template: str,
        variables: Dict[str, Any],
        provider: str = "openai",
        model: str = None
    ) -> ExecutionResult:
        provider_service = self.providers[provider]
        return await provider_service.complete(template, model, **variables)

```

### Custom Evaluation Functions:

```
# Future: Allow users to upload custom evaluation code
```

```

class CustomEvaluator:
    def __init__(self, code: str):
        """Compile user-provided evaluation function."""
        # Sandboxed execution environment
        self.evaluator = self._compile_safely(code)

    def evaluate(self, actual: str, expected: str, metadata: dict) -> float:
        """Execute custom evaluation function."""
        # Run in sandboxed environment
        return self.evaluator(actual, expected, metadata)

```

```

# User uploads:
def evaluate(actual, expected, metadata):
    # Custom logic: check for keywords, regex, etc.
    if "error" in actual.lower():
        return 0.0
    # ... custom scoring
    return similarity_score(actual, expected)

```

Advanced Analytics:

```

# Future: Advanced analytics and insights
class AnalyticsEngine:
    def analyze_trends(self, system_id: UUID) -> TrendAnalysis:
        """Analyze performance trends and predict issues."""
        # Time series analysis
        # Anomaly detection
        # Predictive modeling
        ...

    def generate_insights(self, system_id: UUID) -> List[Insight]:
        """Generate actionable insights."""
        insights = []

        # Cost optimization insights
        if self._can_reduce_cost_without_quality_loss():
            insights.append(Insight(
                type="cost_optimization",
                message="Consider using gpt-4o-mini for 50% cost savings with minimal quality impact",
                impact="high"
            ))

        # Quality insights
        if self._detect_degradation():
            insights.append(Insight(
                type="quality_alert",
                message="Performance has declined 15% over the past week",
            ))

```

```

        impact="critical"
    ))
}

return insights

```

### Integration APIs:

```

# Future: Webhook integrations
@app.post("/api/webhooks")
async def create_webhook(webhook: WebhookCreate):
    """Create webhook for test run events."""
    # Store webhook configuration
    ...

# Trigger webhooks on events
async def trigger_webhook(event: str, data: dict):
    """Send webhook notification."""
    webhooks = await get_webhooks_for_event(event)
    for webhook in webhooks:
        await httpx.post(webhook.url, json=data)

# Slack integration example
async def notify_slack(test_run: TestRun):
    """Send Slack notification."""
    message = f"Test run completed: {test_run.avg_score:.2f} average score"
    await slack_client.post_message(channel="#ai-monitoring", text=message)

```

### Production Readiness

#### Authentication and Authorization:

```

# Future: Add authentication
from fastapi.security import OAuth2PasswordBearer
from jose import JWTError, jwt

oauth2_scheme = OAuth2PasswordBearer(tokenUrl="token")

async def get_current_user(token: str = Depends(oauth2_scheme)):
    """Verify JWT token and return user."""
    credentials_exception = HTTPException(
        status_code=401,
        detail="Could not validate credentials"
    )
    try:
        payload = jwt.decode(token, SECRET_KEY, algorithms=[ALGORITHM])
        user_id: str = payload.get("sub")
        if user_id is None:

```

```

        raise credentials_exception
    except JWTError:
        raise credentials_exception

    user = await get_user(user_id)
    return user

# Protect endpoints
@app.post("/api/prompt-systems")
async def create_prompt_system(
    system: PromptSystemCreate,
    current_user: User = Depends(get_current_user)
):
    # User can only create systems for their organization
    ...

```

#### Audit Logging and Compliance:

```

# Future: Comprehensive audit logging
class AuditLogger:
    async def log_action(
        self,
        user_id: UUID,
        action: str,
        resource_type: str,
        resource_id: UUID,
        metadata: dict = None
    ):
        """Log all user actions for compliance."""
        audit_log = AuditLog(
            user_id=user_id,
            action=action, # "create", "update", "delete", "execute"
            resource_type=resource_type, # "prompt_system", "test_run"
            resource_id=resource_id,
            metadata=metadata,
            ip_address=request.client.host,
            timestamp=datetime.utcnow()
        )
        await db.add(audit_log)
        await db.commit()

```

#### Disaster Recovery:

```

# Backup strategy
# 1. Database backups (automated daily)
# 2. Point-in-time recovery (WAL archiving)
# 3. Test data export (JSON backups)

```

```

async def export_system_data(system_id: UUID) -> dict:
    """Export all data for a system (backup)."""
    system = await get_prompt_system(system_id)
    regression_sets = await get_regression_sets(system_id)
    test_runs = await get_test_runs(system_id)

    return {
        "system": system.dict(),
        "regression_sets": [rs.dict() for rs in regression_sets],
        "test_runs": [tr.dict() for tr in test_runs]
    }

async def import_system_data(data: dict) -> UUID:
    """Import system data (restore)."""
    # Create system from backup
    ...

```

### Performance Monitoring:

```

# Future: APM integration (DataDog, New Relic, etc.)
from datadog import initialize, statsd

initialize(api_key=settings.DATADOG_API_KEY)

@app.middleware("http")
async def datadog_middleware(request: Request, call_next):
    start_time = time.time()
    response = await call_next(request)
    duration = time.time() - start_time

    statsd.histogram(
        'api.request.duration',
        duration,
        tags=[
            f"method:{request.method}",
            f"endpoint:{request.url.path}",
            f"status:{response.status_code}"
        ]
    )

    return response

```

---

## Appendices

### A. Complete API Reference

Base URL: <http://localhost:8000/api>

**Prompt Systems** **GET /prompt-systems** - List all prompt systems - Response: List[PromptSystemResponse]

**POST /prompt-systems** - Create new prompt system - Request: PromptSystemCreate - Response: PromptSystemResponse (201)

**GET /prompt-systems/{id}** - Get prompt system by ID - Response: PromptSystemResponse

**PUT /prompt-systems/{id}** - Update prompt system - Request: PromptSystemUpdate - Response: PromptSystemResponse

**DELETE /prompt-systems/{id}** - Delete prompt system (cascades to regression sets) - Response: 204 No Content

**Regression Sets** **GET /prompt-systems/{system\_id}/regression-sets** - List regression sets for a system - Response: List[RegressionSetResponse]

**POST /prompt-systems/{system\_id}/regression-sets** - Create regression set - Request: RegressionSetCreate - Response: RegressionSetResponse (201)

**DELETE /prompt-systems/{system\_id}/regression-sets/{regression\_id}** - Delete regression set - Response: 204 No Content

**Test Runs** **POST /test-runs** - Create and start test run - Request: TestRunCreate - Response: {"test\_run\_id": "uuid", "status": "pending"} (202 Accepted)

**GET /test-runs** - List test runs - Query params: system\_id (optional), limit (default: 50) - Response: List[TestRunResponse]

**GET /test-runs/{test\_run\_id}** - Get test run status - Response: TestRunResponse

**GET /test-runs/{test\_run\_id}/results** - Get detailed test results - Response: List[TestResultResponse]

**Schedules** **GET /schedules** - List schedules - Query params: system\_id (optional) - Response: List[ScheduleResponse]

**POST /schedules** - Create schedule - Request: ScheduleCreate - Response: ScheduleResponse (201)

**PUT /schedules/{schedule\_id}** - Update schedule - Request: ScheduleUpdate  
- Response: ScheduleResponse

**DELETE /schedules/{schedule\_id}** - Delete schedule - Response: 204 No Content

**POST /schedules/{schedule\_id}/trigger** - Manually trigger scheduled test - Response: {"message": "Test run triggered", "test\_run\_id": "uuid"}

## B. Database Schema

### Complete ER Diagram:

prompt\_systems (1) < (many) regression\_sets

(1)

< (many) test\_runs

(1)

< (many) test\_results

(many)

> (1) regression\_sets

prompt\_systems (1) < (many) schedules

**Table Relationships:** - prompt\_systems → regression\_sets (1:many, CASCADE DELETE)  
- prompt\_systems → test\_runs (1:many) - prompt\_systems → schedules (1:many, CASCADE DELETE) - test\_runs → test\_results (1:many, CASCADE DELETE) - regression\_sets → test\_results (1:many)

**Index Definitions:** - idx\_test\_runs\_system\_id on test\_runs(system\_id)  
- idx\_test\_results\_test\_run\_id on test\_results(test\_run\_id) -  
idx\_prompt\_systems\_last\_reset\_date on prompt\_systems(last\_reset\_date)

## C. Configuration Guide

### Environment Variables:

```
# Database
DATABASE_URL=postgresql+asyncpg://user:pass@host:5432/dbname
```

```
# OpenAI
OPENAI_API_KEY=sk-proj-xxxxxxxxxxxxxxxxxxxxxx
```

```

# Server
BACKEND_HOST=0.0.0.0
BACKEND_PORT=8000
FRONTEND_URL=http://localhost:5173

# CORS
ALLOWED_ORIGINS=http://localhost:5173,http://localhost:3000

```

```

# Optional: Feature flags
ENABLE_SCHEDULING=true
ENABLE_BUDGET_CONTROLS=true
MAX_CONCURRENT_TESTS=10

```

#### Feature Flags:

```

# backend/app/core/config.py
class Settings(BaseSettings):
    ENABLE_SCHEDULING: bool = True
    ENABLE_BUDGET_CONTROLS: bool = True
    MAX_CONCURRENT_TESTS: int = 10
    LOG_LEVEL: str = "INFO"

```

#### Performance Tuning:

```

# Database connection pool
pool_size=20          # Base connections
max_overflow=40        # Additional under load
pool_recycle=3600      # Recycle after 1 hour

# Async execution
MAX_CONCURRENT_TESTS=10 # Limit parallel test execution

# Caching (if Redis added)
REDIS_URL=redis://localhost:6379
CACHE_TTL=300 # 5 minutes

```

## D. Troubleshooting Guide

### Common Issues and Solutions:

- 1. Database Connection Issues - Error:** Connection refused or Unable to connect - **Solutions:** - Verify database is running - Check connection string format (`postgresql+asyncpg://...`) - Verify credentials - Check firewall/network settings
- 2. OpenAI API Errors - Error:** `AuthenticationError` or `Insufficient quota` - **Solutions:** - Verify API key is correct - Check account has credits - Verify key permissions - Check rate limits

**3. Tests Stuck in “Pending” - Error:** Tests never start executing - **Solutions:** - Check background task logs - Verify database connection in background task - Check for exceptions in logs - Restart backend server

**4. Budget Exceeded Errors - Error:** BudgetExceededException - **Solutions:** - Increase `max_daily_cost_usd` for system - Wait for daily reset (midnight UTC) - Manually reset: `UPDATE prompt_systems SET daily_cost_usd = 0, is_budget_paused = false`

**5. Scheduler Not Running - Error:** Scheduled tests don’t execute - **Solutions:** - Check scheduler is started: `GET /health` (check `scheduler` field) - Verify schedule is active: `active = true` - Check cron expression is valid - Review scheduler logs

#### Debug Procedures:

```
# Enable debug logging
import logging
logging.basicConfig(level=logging.DEBUG)

# Check database connection
async def test_db_connection():
    async with AsyncSessionLocal() as db:
        result = await db.execute(text("SELECT 1"))
        print("Database connected:", result.scalar())

# Check OpenAI connection
async def test_openai():
    client = AsyncOpenAI(api_key=settings.OPENAI_API_KEY)
    try:
        response = await client.models.list()
        print("OpenAI connected:", len(response.data), "models")
    except Exception as e:
        print("OpenAI error:", e)
```

#### Performance Optimization:

```
-- Analyze query performance
EXPLAIN ANALYZE SELECT * FROM test_runs WHERE system_id = '...';

-- Check index usage
SELECT * FROM pg_stat_user_indexes WHERE schemaname = 'public';

-- Vacuum and analyze
VACUUM ANALYZE test_runs;
```

---

## Conclusion

This technical specification provides comprehensive documentation of the AI Prompt System Monitor application, covering:

- **Architecture & Design:** Detailed system architecture, component design, data flows
- **Implementation:** Core business logic, evaluation algorithms, advanced features
- **Technical Deep Dives:** Database design, API design, security implementation
- **Performance & Scalability:** Benchmarks, optimization strategies, scaling approaches
- **Testing & QA:** Unit, integration, and E2E testing strategies
- **Deployment & Operations:** Infrastructure, CI/CD, monitoring
- **Challenges & Solutions:** Real-world problems and their solutions
- **Future Roadmap:** Planned enhancements and production readiness features

This document serves as a complete reference for:  
- **Technical Interviews:** Deep technical questions about design choices and trade-offs  
- **Onboarding:** New developers understanding the system  
- **Maintenance:** Future enhancements and bug fixes  
- **RAG Systems:** Rich context for intelligent retrieval and question answering

The application demonstrates production-grade engineering practices including:  
- Async-first architecture for scalability - Concurrent-safe operations (row-level locking)  
- Comprehensive error handling and recovery - Cost management and budget controls  
- Real-time progress tracking - Automated monitoring and alerting  
- Multi-model experimentation and comparison

**Total Document Length:** ~50-70 pages when formatted as PDF

**Code Examples:** 100+ code snippets

**Architecture Diagrams:** 5+ diagrams

**API Endpoints:** 15+ documented endpoints

---

**End of Technical Specification**